Searching in Metric Spaces

EDGAR CHÁVEZ

Escuela de Ciencias Físico-Matemáticas, *Universidad Michoacana*

GONZALO NAVARRO, RICARDO BAEZA-YATES

Depto. de Ciencias de la Computación, *Universidad de Chile*

AND

JOSÉ LUIS MARROQUÍN

Centro de Investigación en Matemáticas (*CIMAT*)

The problem of searching the elements of a set that are close to a given query element under some similarity criterion has a vast number of applications in many branches of computer science, from pattern recognition to textual and multimedia information retrieval. We are interested in the rather general case where the similarity criterion defines a *metric space*, instead of the more restricted case of a vector space. Many solutions have been proposed in different areas, in many cases without cross-knowledge. Because of this, the same ideas have been reconceived several times, and very different presentations have been given for the same approaches. We present some basic results that explain the intrinsic difficulty of the search problem. This includes a quantitative definition of the elusive concept of "intrinsic dimensionality." We also present a unified

Categories and Subject Descriptors: F.2.2 [**Analysis of algorithms and problem complexity**]: Nonnumerical algorithms and problems—*computations on discrete structures*, *geometrical problems and computations*, *sorting and searching*; H.2.1 [**Database management**]: Physical design—*access methods*; H.3.1 [**Information storage and retrieval**]: Content analysis and indexing—*indexing methods*; H.3.2 [**Information storage and retrieval**]: Information storage—*file organization*; H.3.3 [**Information storage and retrieval**]: Information search and retrieval—*clustering*, *search process*; I.5.1 [**Pattern recognition**]: Models—*geometric*; I.5.3 [**Pattern recognition**]: Clustering

General Terms: Algorithms

Additional Key Words and Phrases: Curse of dimensionality, nearest neighbors, similarity searching, vector spaces

---

view of all the known proposals to organize metric spaces, so as to be able to understand them under a common framework. Most approaches turn out to be variations on a few different concepts. We organize those works in a taxonomy that allows us to devise new algorithms from combinations of concepts not noticed before because of the lack of communication between different communities. We present experiments validating our results and comparing the existing approaches. We finish with recommendations for practitioners and open questions for future development.

1. INTRODUCTION

Searching is a fundamental problem in computer science, present in virtually every computer application. Simple applications pose simple search problems, whereas a more complex application will require, in general, a more sophisticated form of searching.

The search operation traditionally has been applied to "structured data," that is, numerical or alphabetical information that is searched for exactly. That is, a search query is given and the number or string that is *exactly equal* to the search query is retrieved. Traditional databases are built around the concept of exact searching: the database is divided into records, each record having a fully comparable key. Queries to the database return all the records whose keys match the search key. More sophisticated searches such as range queries on numerical keys or prefix searching on alphabetical keys still rely on the concept that two keys are or are not equal, or that there is a total linear order on the keys. Even in recent years, when databases have included the ability to store new data types such as images, the search has still been done on a predetermined number of keys of numerical or alphabetical types.

With the evolution of information and communication technologies, unstructured repositories of information have emerged. Not only new data types such as free text, images, audio, and video have to be queried, but also it is no longer possible to structure the information in keys and records. Such structuring is very difficult (either manually or computationally) and restricts beforehand the types of queries that can be posed later. Even when a classical structuring is possible, new applications such as data mining require accessing the database by any field, not only those marked as "keys." Hence, new models for searching in unstructured repositories are needed.

The above scenarios require more general search algorithms and models than those classically used for simple data. A unifying concept is that of "similarity searching" or "proximity searching," that is, searching for database elements that are similar or close to a given query element.[1] Similarity is modeled with a distance function that satisfies the triangle inequality, and the set of objects is called a *metric space*. Since the problem has appeared in many diverse areas, solutions have appeared in many unrelated fields, such as statistics, computational geometry, artificial intelligence, databases, computational biology, pattern recognition, and data mining, to name a few. Since the current solutions come from such diverse fields, it is not surprising that the same solutions have been reinvented many times, that obvious combinations of solutions have not been noticed, and that no thorough comparisons have been done. More important, there have been no attempts to conceptually unify all those solutions.

In some applications the metric space turns out to be of a particular type called "vector space," where the elements consist of k real-valued coordinates. A lot of work has been done on vector spaces by exploiting their geometric properties, but normally these cannot be extended to general metric spaces where the only available

---

[1] The term "approximate searching" is also used, but it is misleading and we use it here only when referring to approximation algorithms.

information is the distance among objects. In this general case, moreover, the distance is normally quite expensive to compute, so the general goal is to reduce the number of distance evaluations. In contrast, the operations in vector spaces tend to be simple and hence the goal is mainly to reduce I/O. Some important advances have been done for general metric spaces, mainly around the concept of building an *index*, that is, a data structure to reduce the number of distance evaluations at query time. Some recent work [Ciaccia et al. 1997; Prabhakar et al. 1998] tries to achieve simultaneously the goals of reducing the number of distance evaluations and the amount of I/O performed.

The main goal of this work is to present a unifying framework to describe and analyze all the existing solutions to this problem. We show that all the existing indexing algorithms for proximity searching consist of building a set of equivalence classes, discarding some classes, and exhaustively searching the rest. Two main techniques based on equivalence relations, namely, pivoting and compact partitions, are shown to encompass all the existing methods. As a consequence of the analysis we are able to build a taxonomy on the existing algorithms for proximity search, to classify them according to their essential features, and to analyze their efficiency. We are able to identify essentially similar approaches, to point out combinations of ideas that have not previously been noticed, and to identify the main open problems in this area. We also present quantitative methods to assert the intrinsic difficulty in searching on a given metric space and provide lower bounds on the search problem. This includes a quantitative definition of the previously conceptual notion of "intrinsic dimensionality," which we show to be very appropriate. We present some experimental results that help to validate our assertions.

We remark that we are concerned with the *essential* features of the search algorithms for general metric spaces. That is, we try to extract the basic features from the wealth of existing solutions, so as to be able to categorize and analyze them under a common framework. We focus mainly on the number of distance evaluations needed to execute *range queries* (i.e., with fixed tolerance radius), which are the most basic ones. However, we also pay some attention to the total CPU time, time and space cost to build the indexes, nearest neighbor queries, dynamic capabilities of the indexes, and I/O considerations. There are some features that we definitely do not cover in order to keep our scope reasonably bounded, such as (1) complex similarity queries involving more than one similarity predicate [Ciaccia et al. 1998b], as few works on them exist and they are an elaboration over the simple similarity queries (a particular case is polygon searching in vector spaces); (2) subqueries (i.e., searching a small element inside a larger element) since the solutions are basically the same after a domain-dependent transformation is done; and (3) inverse queries (i.e., finding the elements for which q is their closest neighbor) and total queries (e.g., finding all the closest neighbors) since the algorithms are, again, built over the simple ones.

This article is organized as follows. The first part (Sections 2 through 5) is a pure survey of the state of the art in searching metric spaces, with no attempt to provide a new way of thinking about the problem. The second part (Sections 6 through 8) presents our basic results on the difficulty of the search problem and our unifying model that allows understanding the essential features of the problem and its existing solutions. Finally, Section 9 gives our conclusions and points out future research directions.

2. MOTIVATING APPLICATIONS

We now present a sample of applications where the concept of proximity searching appears. Since we have not presented a formal model yet, we do not try to explain the connections between the different applications. We rather delay this discussion to Section 3.

2.1. Query by Content in Structured Databases

In general, the query posed to a database presents a *piece* of a record of information, and it needs to retrieve the *entire* record. In the classical approach, the piece presented is fixed (the *key*). Moreover, it is not allowed to search with an incomplete or an erroneous key. On the other hand, in the more general approach required nowadays the concept of searching with a key is generalized to searching with an arbitrary subset of the record, allowing errors or not.

Possible types of searches are *point* or *key* search (all the key information is given), *range* search (only some fields are given or only a range of values is specified for them), and *proximity* search (in addition, records "close" to the query are considered interesting). These types of search are of use in data mining (where the interesting parts of the record cannot be predetermined), when the information is not precise, when we are looking for a range of values, when the search key may have errors (e.g., a misspelled word), and so on.

A general solution to the problem of range queries by any record field is the *grid file* [Nievergelt and Hinterberger 1984]. The domain of the database is seen as a hyperrectangle of k dimensions (one per record field), where each dimension has an ordering according to the domain of the field (numerical or alphabetical). Each record present in the database is considered as a point inside the hyperrectangle. A query specifies a subrectangle (i.e., a range along each dimension), and all the points inside the specified query are retrieved. This does not address the problem of searching on nontraditional data types, nor allowing errors that cannot be recovered with a range query. However, it converts the original search problem to a problem of obtaining, in a given space, all the points "close" to a given query point. Grid files are essentially a disk organization technique to efficiently retrieve range queries in secondary memory.

2.2. Query by Content in Multimedia Objects

New data types such as images, fingerprints, audio, and video (called "multimedia" data types) cannot be meaningfully queried in the classical sense. Not only can they not be ordered, but they cannot even be compared for equality. There is no interest in an application for searching an audio segment exactly equal to a given one. The probability that two different images are pixelwise equal is negligible unless they are digital copies of the same source. In multimedia applications, all the queries ask for objects *similar* to a given one. Some example applications are face recognition, fingerprint matching, voice recognition, and, in general, multimedia databases [Apers et al. 1997; Yoshitaka and Ichikawa 1999].

Think, for example, of a repository of images. Interesting queries are of the type, "*Find an image of a lion with a savanna background.*" If the repository is tagged, and each tag contains a full description of what is inside the image, then our example query can be solved with a classical scheme. Unfortunately, such a classification cannot be done automatically with the available image processing technology. The state of object recognition in real-world scenes is still too immature to perform such complex tasks. Moreover, we cannot predict all the possible queries that will be posed so as to tag the image for every possible query. An alternative to automatic classification consists of considering the query as an *example* image, so that the system searches all the images similar to the query. This can be built inside a more complex feedback system where the user approves or rejects the images found, and a new query is submitted with the approved images. It is also possible that the query is just part of an image and the system has to retrieve the whole image.

These approaches are based on the definition of a similarity function among objects. Those functions are provided by an expert, but they pose no assumptions on the type of queries that can be answered. In many cases, the distance is obtained via

a set of k "features" that are extracted from the object (e.g., in an image a useful feature is the average color). Then each object is represented as its k features (i.e., a point in a k-dimensional space), and we are again in a case of range queries on vector spaces.

There is a growing community of scientists deeply involved with the development of such similarity measures [Cascia et al. 1998; Bhanu et al. 1998; Bimbo and Vicario 1998].

2.3. Text Retrieval

Although not considered a multimedia data type, unstructured text retrieval poses similar problems to those of multimedia retrieval. This is because textual documents are in general not structured to easily provide the desired information. Text documents may be searched for strings that are present or not, but in many cases they are searched for semantic concepts of interest. For instance, an ideal scenario would allow searching a text dictionary for a concept such as, "*to free from obligation*," retrieving the word "*redeem*." This search problem cannot be properly stated with classical tools.

A large community of researchers has been working on this problem for a long time [Salton and McGill 1983; Frakes and Baeza-Yates 1992; Baeza-Yates and Ribeiro-Neto 1999]. A number of measures of similarity have emerged. The problem is solved basically by retrieving documents similar to a given query. The user can even present a document as a query, so that the system finds similar documents. Some similarity approaches are based on mapping a document to a vector of real values, so that each dimension is a vocabulary word and the relevance of the word to the document (computed using some formula) is the coordinate of the document along that dimension. Similarity functions are then defined in that space. Notice, however, that the dimensionality of the space is very high (thousands of dimensions).

Another problem related to text retrieval is spelling. Since huge text databases with low quality control are emerging (e.g., the Web), and typing, spelling, or OCR (optical character recognition) errors are commonplace in the text and the query, we have cases where documents that contain a misspelled word are no longer retrievable by a correctly written query. Models of similarity among words exist (variants of the "edit distance" [Sankoff and Kruskal 1983]) that capture very well those types of errors. In this case, we give a word and want to retrieve all the words close to it. Another related application is spelling checkers, where we look for close variants of the misspelled word.

In particular, OCR can be done using a low-level classifier, so that misspelled words can be corrected using the edit distance to find promising alternatives to replace incorrect words [Chávez and Navarro 2002].

2.4. Computational Biology

ADN and protein sequences are the basic object of study in molecular biology. As they can be modeled as texts, we have the problem of finding a given sequence of characters inside a longer sequence. However, an exact match is unlikely to occur, and computational biologists are more interested in finding parts of a longer sequence that are similar to a given short sequence. The fact that the search is not exact is due to minor differences in the genetic streams that describe beings of the same or closely related species. The measure of similarity used is related to the probability of mutations such as reversals of pieces of the sequences and other rearrangements [Waterman 1995; Sankoff and Kruskal 1983].

Other related problems are to build phylogenetic trees (a tree sketching the evolutionary path of the species), to search patterns for which only some properties are known, and others.

2.5. Pattern Recognition and Function Approximation

A simplified definition of pattern recognition is the construction of a function

approximator. In this formulation of the problem one has a finite sample of the data, and each data sample is labeled as belonging to a certain class. When a fresh data sample is provided, the system is required to label this new sample with one of the known data labels. In other words, the classifier can be thought of as a function defined from the object (data) space to the set of labels. In this sense all the classifiers are considered as function approximators.

If the objects are m-dimensional vectors of real numbers then the natural choices are neural nets and fuzzy function approximators. Another popular universal function approximator, the k-nearest-neighbor classifier, consists of finding the k objects nearest to the unlabeled sample, and assigning to this sample the label having the majority among the k nearest objects. Opposed to neural nets and fuzzy classifiers, the k-nearest-neighbor rule has zero training time, but if no indexing algorithm is used it has linear complexity [Duda and Hart 1973].

Other applications of this universal function approximator are density estimation [Devroye 1987] and reinforcement learning [Sutton and Barto 1998]. In general, any problem where we want to infer a function based on a finite set of samples is a potential application.

2.6. Audio and Video Compression

Audio and video transmission over a narrowband channel is an important problem, for example, in Internet-based audio and video conferencing or in wireless communication. A frame (a static picture in a video, or a fragment of the audio) can be thought of as formed by a number of (possibly overlapped) subframes (16×16 subimages in a video, for example). In a very succinct description, the problem can be solved by sending the first frame as-is and for the next frames sending only the subframes having a significative difference from the previously sent subframes. This description encompasses the MPEG standard.

The algorithms use, in fact, a subframe buffer. Each time a frame is about to be sent it is searched (with a tolerance) in the subframe buffer and if it is not found then the entire subframe is added to the buffer. If the subframe is found then only the index of the similar frame found is sent. This implies, naturally, that a fast similarity search algorithm has to be incorporated to the server to maintain a minimum frames-per-second rate.

3. BASIC CONCEPTS

All the applications presented in the previous section share a common framework, which is in essence to find close objects, under some suitable similarity function, among a finite set of elements. In this section we present the formal model comprising all the above cases.

3.1. Metric Spaces

We now introduce the basic notation for the problem of satisfying proximity queries and for the model used to group and analyze the existing algorithms.

The set $\mathbb{X}$ denotes the universe of *valid* objects. A finite subset of it, $\mathbb{U}$, of size $n = |\mathbb{U}|$, is the set of objects where we search. $\mathbb{U}$ is called the *dictionary*, *database*, or simply our set of *objects* or *elements*. The function

$$d : \mathbb{X} \times \mathbb{X} \to \mathbb{R}$$

denotes a measure of "distance" between objects (i.e., the smaller the distance, the closer or more similar are the objects). Distance functions have the following properties.

$(p1)$ $\forall x, y \in \mathbb{X}, d(x, y) \geq 0$ positiveness,

$(p2)$ $\forall x, y \in \mathbb{X}$,
     $d(x, y) = d(y, x)$ symmetry,

$(p3)$ $\forall x \in \mathbb{X}, d(x, x) = 0$ reflexivity,

and in most cases

$(p4)$ $\forall x, y \in \mathbb{X}, x \neq y \Rightarrow$
     $d(x, y) > 0$ strict
                          positiveness.

The similarity properties enumerated above only ensure a consistent definition of the function, and cannot be used to save

comparisons in a proximity query. If d is indeed a *metric*, that is, if it satisfies

$(p5)$ $\forall x, y, z \in \mathbb{X}, \; d(x, y) \leq d(x, z) + d(z, y)$ triangle inequality,

then the pair $(\mathbb{X}, d)$ is called a *metric space*.

If the distance does not satisfy the strict positiveness property $(p4)$ then the space is called a *pseudometric space*. Although for simplicity we do not consider pseudo-metric spaces in this work, all the presented techniques are easily adapted to them by simply identifying all the objects at distance zero as a single object. This works because, if $(p5)$ holds, one can easily prove that $d(x, y) = 0 \Rightarrow \forall z, d(x, z) = d(y, z)$.

In some cases we may have a *quasi-metric*, where the symmetry property $(p2)$ does not hold. For instance, if the objects are corners in a city and the distance corresponds to how much a car must travel to move from one to the other, then the existence of one-way streets makes the distance asymmetric. There exist techniques to derive a new, symmetric, distance function from an asymmetric one, such as $d'(x, y) = d(x, y) + d(y, x)$. However, to be able to bound the search radius of a query when using the symmetric function we need specific knowledge of the domain.

Finally, we can relax the triangle inequality $(p5)$ to $d(x, y) \leq \alpha d(x, z) + \beta d(z, y) + \delta$, and after some scaling we can search in this space using the same algorithms designed for metric spaces. If the distance is symmetric we need $\alpha = \beta$ for consistency.

In the rest of the article we use the term *distance* in the understanding that we refer to a metric.

3.2. Proximity Queries

There are basically three types of queries of interest in metric spaces.

Range query $(q, r)_d$. Retrieve all elements that are within distance r to q. This is $\{u \in \mathbb{U}/d(q, u) \leq r\}$.

Nearest neighbor query $NN(q)$. Retrieve the closest elements to q in $\mathbb{U}$. This is $\{u \in \mathbb{U}/\forall v \in \mathbb{U}, \; d(q, u) \leq d(q, v)\}$. In some cases we are satisfied with one such element (in continuous spaces there is normally just one answer anyway). We can also give a maximum distance r^* such that if the closest element is at distance more than r^* we do not want any one reported.

k-Nearest neighbor query $NN_k(q)$. Retrieve the k closest elements to q in $\mathbb{U}$. This is, retrieve a set $A \subseteq \mathbb{U}$ such that $|A| = k$ and $\forall u \in A, v \in \mathbb{U} - A$, $d(q, u) \leq d(q, v)$. Note that in case of ties we are satisfied with any set of k elements satisfying the condition.

The most basic type of query is the range query. The left part of Figure 1 illustrates a query on a set of points which is our running example, using $\mathbb{R}^2$ as the metric space for clarity.
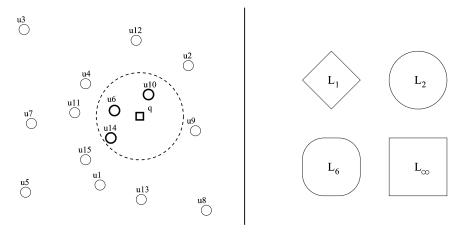
A *range query* is therefore a pair $(q, r)_d$ with q an element in $\mathbb{X}$ and r a real number indicating the *radius* (or tolerance) of the query. The set $\{u \in \mathbb{U}, \; d(q, u) \leq r\}$ is called the *outcome* of the range query.

We use "NN" as an abbreviation of "nearest neighbor," and give the generic name "NN-query" to the last two types of queries and "NN searching" to the techniques to solve them. As we see later, NN-queries can be systematically built over range queries.

The total time to evaluate a query can be split as

$$
\begin{aligned}
T \;=\; & \text{distance evaluations} \\
& \times \text{complexity of } d() \\
& + \text{extra CPU time} + \text{I/O time}
\end{aligned}
$$

and we would like to minimize T. In many applications, however, evaluating $d()$ is so costly that the other components of the cost can be neglected. This is the model we use in this article, and hence the *number* of distance evaluations performed is the measure of the complexity of the algorithms. We can even allow a linear in n (but reasonable) amount of CPU work and a linear traversal over the database on disk, as long as the number of distance computations is kept low. However, we pay some marginal attention to

Fig. 1.   On the left, an example of a range query on a set of points. On the right, the set of points at the same distance to a center point, for different L_s distances.

the so-called *extra CPU time*. The I/O time can be the dominant factor in some applications and negligible in others, depending on the amount of main memory available and the relative cost to compute the distance function. We cover the little existing work on relating metric spaces and I/O considerations in Section 5.3.2.

It is clear that either type of query can be answered by examining the entire dictionary $\mathbb{U}$. In fact if we are not allowed to preprocess the data (i.e., to build an index data structure), then this exhaustive examination is the only way to proceed. An *indexing algorithm* is an offline procedure to build beforehand a data structure (called an *index*) designed to save distance computations when answering proximity queries later. This data structure can be expensive to build, but this will be amortized by saving distance evaluations over many queries to the database. The aim is therefore to design efficient indexing algorithms to reduce the number of distance evaluations. All these structures work on the basis of discarding elements using the triangle inequality (the only property that allows saving distance evaluations).

4. THE CASE OF VECTOR SPACES

If the elements of the metric space $(\mathbb{X}, d)$ are indeed tuples of real numbers (actu-

ally tuples of any field) then the pair is called a *finite-dimensional vector space*, or vector space for short.

A k-dimensional vector space is a particular metric space where the objects are identified with k real-valued coordinates $(x_1, \ldots, x_k)$. There are a number of options for the distance function to use, but the most widely used is the family of L_s distances, defined as

$$L_s((x_1, \ldots, x_k), (y_1, \ldots, y_k))$$
$$= \left(\sum_{i=1}^{k} |x_i - y_i|^s \right)^{1/s}.$$

The right part of Figure 1 illustrates some of these distances. For instance, the L_1 distance accounts for the sum of the differences along the coordinates. It is also called "block" or "Manhattan" distance, since in two dimensions it corresponds to the distance to walk between two points in a city of rectangular blocks. The L_2 distance is better known as "Euclidean" distance, as it corresponds to our notion of spatial distance. The other most used member of the family is L_∞, which corresponds to taking the limit of the L_s formula when s goes to infinity. The result is that the distance between two points is the

maximum difference along a coordinate:

$$L_\infty((x_1, \ldots, x_k), (y_1, \ldots, y_k))$$
$$= \max_{i=1}^{k} |x_i - y_i|.$$

Searching with the L_∞ distance corresponds directly to a classical range search query, where the range is the k-dimensional hyperrectangle. This distance plays a special role in this survey.

In many applications the metric space is indeed a vector space; that is, the objects are k-dimensional points and the similarity is interpreted geometrically. A vector space permits more freedom than a general metric space when designing search approaches, since it is possible to use geometric and coordinate information that is unavailable in a general metric space.

In this framework optimal algorithms (on the database size) exist in both the average and the worst case [Bentley et al. 1980] for closest point search. Search structures for vector spaces are called *spatial access methods* (SAM). Among the most popular are kd-trees [Bentley 1975, 1979], R-trees [Guttman 1984], quadtrees [Samet 1984], and the more recent X-trees [Berchtold et al. 1996]. These techniques make extensive use of coordinate information to group and classify points in the space. For example, kd-trees divide the space along different coordinates and R-trees group points in hyperrectangles. Unfortunately the existing techniques are very sensitive to the vector space dimension. Closest point and range search algorithms have an exponential dependency on the dimension of the space [Chazelle 1994] (this is called the *curse of dimensionality*).

Vector spaces may suffer from large differences between their *representational* dimension (k) and their *intrinsic* dimension (i.e., the real number of dimensions in which the points can be embedded while keeping the distances among them). For example, a plane embedded in a 50-dimensional space has intrinsic dimension 2 and representational dimension 50. This is, in general, the case of real applications, where the data are clustered, and it has

led to attempts to measure the intrinsic dimension such as the concept of "fractal dimension" [Faloutsos and Kamel 1994]. Despite the fact that no technique can cope with intrinsic dimension higher than 20, much higher representational dimensions can be handled by dimensionality reduction techniques [Faloutsos and Lin 1995; Cox and Cox 1994; Hair et al. 1995].

Since efficient techniques to cope with vector spaces exist, application designers try to give their problems a vector space structure. However, this is not always easy or feasible at all. For example, experts in image processing try to express the similarity between images as the distance between "vectors" of features extracted from the images, although in many cases better results are obtained by coding specific functions that compare two images, despite their inability to be easily expressed as the distance between two vectors (e.g., crosstalk between features [Faloutsos et al. 1994]). Another example that resists conversion into a vector space is similarity functions between strings, to compare DNA sequences, for instance. A step towards improving this situation is Chávez and Navarro [2002].

For this reason several authors resort to general metric spaces, even knowing that the search problem is much more difficult. Of course it is also possible to treat a vector space as a general metric space, by using only the distances between points. One immediate advantage is that the intrinsic dimension of the space shows up, independent of any representational dimension (this requires extra care in vector spaces). It is interesting to remark that Ciaccia et al. [1997] present preliminary results showing that a metric space data structure (the M-tree) can outperform a well-known vector space data structure (the R*-tree) when applied to a vector space.

Specific techniques for vector spaces are a whole different world which we do not intend to cover in this work (see Samet [1984], White and Jain [1996], Gaede and Günther [1998], and Böhm et al. [2002] for good surveys). However, we discuss in the next section a technique that, instead of treating a vector space as a metric space,

tries to embed a general metric space into a vector space. This concept is central in this survey, although specific knowledge of specific techniques for vector spaces is, as we shortly show, not necessary to understand it.

4.1. Resorting to Vector Spaces

An interesting and natural reduction of the similarity search problem consists of a mapping Φ from the original metric space into a vector space. In this way, each element of the original metric space is represented as a point in the target vector space. The two spaces are related by two distances: the original one $d(x, y)$ and the distance in the projected space $D(\Phi(x), \Phi(y))$. If the mapping is *contractive* (i.e., $D(\Phi(x), \Phi(y)) \leq d(x, y)$ for any pair of elements), then one can process range queries in the projected space with the same radius. Since some spurious elements can be captured in the target space, the outcome of the query in the projected space is a *candidate list*, which is later verified elementwise with the original distance to obtain the actual outcome of the query.

Intuitively, the idea is to "invent" k coordinates and map the points onto a vector space, using some vector space technique as a first filter to the actual answer to a query. One of the main theses of this work is that a large subclass of the existing algorithms can be regarded as relying on some mapping of this kind. A widely used method (explained in detail in Section 6.6) is to select $\{p_1, \ldots, p_k\} \subseteq \mathbb{U}$ and map each $u \in \mathbb{U}$ to $(\mathbb{R}^k, L_\infty)$ using $\Phi(u) = (d(u, p_1), \ldots, d(u, p_k))$. It can be seen that this mapping is contractive but not proximity preserving.

If, on the other hand, the mapping is *proximity preserving* (i.e., $d(x, y) \leq d(x, z) \Rightarrow D(\Phi(x), \Phi(y)) \leq D(\Phi(x), \Phi(z))$), then NN-queries can be directly performed in the projected space. Indeed, most current algorithms for NN-queries are based in range queries, and with some care they can be done in the projected space if the mapping is contractive, even if it is not proximity preserving.

This type of mapping is a special case of a general idea in the literature that says that one can find it cheaper to compute distances that lower-bound the real one, and use the cheaper distance to filter out most elements (e.g., for images, the average color is cheaper to compute than the differences in the color histograms). While in general this is domain-dependent, mapping onto a vector space can be done without knowledge of the domain. After the mapping is done and we have identified each data element with a point in the projected space, we can use a general-purpose spatial access method for vector spaces to retrieve the candidate list. The elements found in the projected space must be finally checked using the original distance function.

Therefore, there are two types of distance evaluations: first to obtain the coordinates in the projected space and later to check the final candidates. These are called "internal" and "external" evaluations, respectively, later in this work. Clearly, incrementing internal evaluations improves the quality of the filter and reduces external evaluations, and therefore we seek a balance.

Notice finally that the search itself in the projected space does not use evaluations of the original distance, and hence it is costless under our complexity measure. Therefore, the use of kd-trees, R-trees, or other data structures aims at reducing the extra CPU time, but it makes no difference in the number of evaluations of the d distance.

How well do metric space techniques perform in comparison to vector space methods? It is difficult to give a formal answer because of the different cost models involved. In metric spaces we use the number of distance evaluations as the basic measure of complexity, while vector space techniques may very well use many coordinate manipulations and not a single evaluation of the distance. Under our model, the cost of a method that maps to a vector space to trim the candidate list is measured as the number of distance evaluations to realize the mapping plus the final distances to filter the trimmed candidate

list, whereas the work on the artificial coordinates is seen as just extra CPU time.

A central question related to this reduction is: how well can a metric space be embedded into a vector space? How many coordinates have to be considered so that the original metric space and the target vector spaces are similar enough so that the candidate list given by the vector space is not much larger than the actual outcome of the query in the original space? This is a very difficult question that lies behind this entire article, and we return to it in Section 7.

The issue is better developed in vector spaces. There are different techniques to reduce the dimensionality of a set of points while preserving the original distances as much as possible [Cox and Cox 1994; Hair et al. 1995; Faloutsos and Lin 1995], that is, to find the intrinsic dimension of the data.

5. CURRENT SOLUTIONS FOR METRIC SPACES

In this section we explain the existing indexes to structure metric spaces and how they are used for range and NN searching. Since we have not yet developed the concepts of a unifying perspective, the description is kept at an intuitive level, without any attempt to analyze why some ideas are better or worse. We add a final subsection devoted to more advanced issues such as dynamic capabilities, I/O considerations, and approximate and probabilistic algorithms.

5.1. Range Searching

We divide the presentation in several parts. The first one deals with tree indexes for *discrete* distance functions, that is, functions that deliver a small set of values. The second part corresponds to tree indexes for *continuous* distance functions, where the set of alternatives is infinite or very large. Third, we consider other methods that are not tree-based.

Table I summarizes the complexities of the different structures. These are obtained from the source papers, which use different (and incompatible) assumptions and in many cases give just gross analyses or no analysis at all (just heuristic considerations). Therefore, we give the complexities as claimed by the authors of each paper, not as a proven fact. At best, the results are analytical but rely on diverse simplifying assumptions. At worst, the results are based on a few incomplete experiments. Keep also in mind that there are hidden factors depending (in many cases exponentially) on the dimension of the space, and that the query complexity is always on average, as in the worst case we can be forced to compare all the elements. Even in the simple case of orthogonal range searching on vector spaces there exist $\Omega(n^\alpha)$ lower bounds for the worst case [Melhorn 1984].

5.1.1. Trees for Discrete Distance Functions. We start by describing tree data structures that apply to distance functions which return a small set of different values. At the end we show how to cope with the general case with these trees.

5.1.1.1. BKT. Probably the first general solution to search in metric spaces was presented in Burkhard and Keller [1973]. They propose a tree (thereafter called the Burkhard–Keller Tree, or BKT), which is suitable for discrete-valued distance functions. It is defined as follows. An arbitrary element $p \in \mathbb{U}$ is selected as the root of the tree. For each distance $i > 0$, we define $\mathbb{U}_i = \{u \in \mathbb{U}, d(u, p) = i\}$ as the set of all the elements at distance i to the root p. Then, for any nonempty $\mathbb{U}_i$, we build a child of p (labeled i), where we recursively build the BKT for $\mathbb{U}_i$. This process can be repeated until there is only one element to process, or until there are no more than b elements (and we store a *bucket* of size b). All the elements selected as roots of subtrees are called *pivots*.

When we are given a query q and a distance r, we begin at the root and enter into all children i such that $d(p, q) - r \le i \le d(p, q) + r$, and proceed recursively. If we arrive at a leaf (bucket of size one or more) we compare sequentially all its elements. Each time we perform a comparison

Table I.　Average Complexities of Existing Approaches[a]

Data Structure	Space Complexity	Construction Complexity	Claimed Query Complexity	Extra CPU Query Time
BKT	n ptrs	$n \log n$	n^α	—
FQT	$n \mathinner{.\,.} n \log n$ ptrs	$n \log n$	n^α	—
FHQT	$n \mathinner{.\,.} nh$ ptrs	nh	$\log n$ [b]	n^α
FQA	nhb bits	nh	$\log n$ [b]	$n^\alpha \log n$
VPT	n ptrs	$n \log n$	$\log n$ [c]	—
MVPT	n ptrs	$n \log n$	$\log n$ [c]	—
VPF	n ptrs	$n^{2-\alpha}$	$n^{1-\alpha} \log n$ [c]	—
BST	n ptrs	$n \log n$	not analyzed	—
GHT	n ptrs	$n \log n$	not analyzed	—
GNAT	nm^2 dsts	$nm \log_m n$	not analyzed	—
VT	n ptrs	$n \log n$	not analyzed	—
MT	n ptrs	$n(m \mathinner{.\,.} m^2) \log_m n$	not analyzed	—
SAT	n ptrs	$n \log n / \log \log n$	$n^{1-\Theta(1/\log \log n)}$	—
AESA	n^2 dsts	n^2	$O(1)$ [d]	$n \mathinner{.\,.} n^2$
LAESA	kn dsts	kn	$k + O(1)$ [d]	$\log n \mathinner{.\,.} kn$
LC	n ptrs	n^2/m	not analyzed	n/m

[a]Time complexity considers only n, not other parameters such as dimension. Space complexity mentions the most expensive storage units used ("ptrs" is short for "pointers" and "dsts" for "distances"). α is a number between 0 and 1, different for each structure, whereas the other letters are parameters particular to each structure.
[b]If $h = \log n$.
[c]Only valid when searching with very small radii.
[d]Empirical conclusions without analysis, in the case of LAESA for "large enough" k.

(against pivots or bucket elements u) where $d(q, u) \le r$, we report the element u.

The triangle inequality ensures that we cannot miss an answer. All the subtrees not traversed contain elements u that are at distance $d(u, p) = i$ from some node p, where $|d(p, q) - i| > r$. By the triangle inequality, $d(p, q) \le d(p, u) + d(u, q)$, and therefore $d(u, q) \ge d(p, q) - d(p, u) > r$.
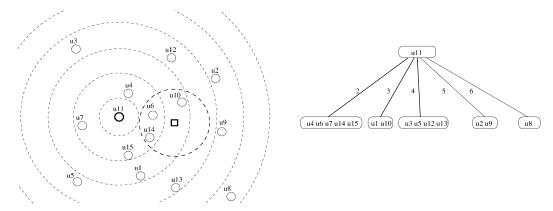
Figure 2 shows an example where the element u_{11} has been selected as the root. We have built only the first level of the BKT for simplicity. A query q is also shown, and we have emphasized the branches of the tree that would have to be traversed. In this and all the examples of this section we discretize the distances of our example, so that they return integer values.

The results of Table I for BKTs are extrapolated from those made for fixed queries trees [Baeza-Yates et al. 1994], which can be easily adapted to this case. The only difference is that the space overhead of BKTs is $O(n)$ because there is exactly one element of the set per tree node.

5.1.1.2. FQT. A further development over BKTs is the fixed queries tree or FQT [Baeza-Yates et al. 1994]. This tree is basi-

cally a BKT where all the pivots stored in the nodes of the same level are the same (and of course do not necessarily belong to the set stored in the subtree). The actual elements are all stored at the leaves. The advantage of such construction is that some comparisons between the query and the nodes are saved along the backtracking that occurs in the tree. If we visit many nodes of the same level, we do not need to perform more than one comparison because all the pivots in that level are the same. This is at the expense of somewhat taller trees. FQTs are shown experimentally in Baeza-Yates et al. [1994] to perform fewer distance evaluations at query time using a couple of different metric space examples. Under some simplifying assumptions (experimentally validated in the paper) they show that FQTs built over n elements are $O(\log n)$ height on average, are built using $O(n \log n)$ distance evaluations, and that the average number of distance computations is $O(n^\alpha)$, where $0 < \alpha < 1$ is a number that depends on the range of the search and the structure of the space (this analysis is easy to extend to BKTs as well). The space complexity is superlinear since, unlike BKTs, it is not

Fig. 2. On the left, the division of the space obtained when u_{11} is taken as a pivot. On the right, the first level of a BKT with u_{11} as root. We also show a query q and the branches that it has to traverse. We have discretized the distances so they return integer values.

true that a different element is placed at each node of the tree. An upper bound is $O(n \log n)$ since the average height is $O(\log n)$.
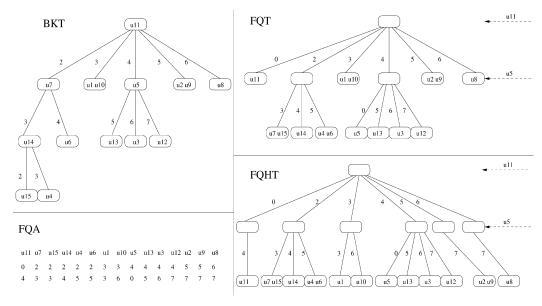
5.1.1.3. FHQT. In Baeza-Yates et al. [1994] and Baeza-Yates [1997], the authors propose a variant called "fixed-height FQT" (or FHQT for short), where all the leaves are at the same depth h, regardless of the bucket size. This makes some leaves deeper than necessary, which makes sense because we may have already performed the comparison between the query and the pivot of an intermediate level, therefore eliminating for free the need to consider the leaf. In Baeza-Yates [1997] and Baeza-Yates and Navarro [1998] it is shown that by using $O(\log n)$ pivots, the search takes $O(\log n)$ distance evaluations (although the cost depends exponentially on the search radius r). The extra CPU time, that is, the number of nodes traversed, remains however, $O(n^{\alpha})$. The space, like FQTs, is somewhere between $O(n)$ and $O(nh)$. In practice the optimal $h = O(\log n)$ cannot be achieved because of space limitations.

5.1.1.4. FQA. In Chávez et al. [2001], the fixed queries array (FQA) is presented. The FQA, although not properly a tree, is no more than a compact representation of the FHQT. Imagine that an FHQT of fixed height h is built on the set. If we tra-

verse the leaves of the tree left to right and put the elements in an array, the result is the FQA. For each element of the array we compute h numbers representing the branches to take in the tree to reach the element from the root (i.e., the distances to the h pivots). Each of these h numbers is coded in b bits and they are concatenated in a single (long) number so that the higher levels of the tree are the most significant digits.

As a result the FQA is sorted by the resulting hb-bits number, each subtree of the FHQT corresponds to an interval in the FQA, and each movement in the FHQT is simulated with two binary searches in the FQA (at $O(\log n)$ extra CPU cost factor, but no extra distances are computed). There is a similarity between this idea and suffix trees versus suffix arrays [Frakes and Baeza-Yates 1992]. This idea of using fewer bits to represent the distances appeared also in the context of vector spaces [Weber et al. 1998].

Using the same memory, the FQA simulation is able to use many more pivots than the original FHQT, which improves the efficiency. The b bits needed by each pivot can be lowered by merging branches of the FHQT, trying that about the same number of elements lies in each cell of the next level. This allows using even more pivots with the same space usage. For reasons that are made clear later, the FQA is also called FMVPA in this work.

Fig. 3. Example BKT, FQT, FHQT, and FQA for our set of points. We use $b = 2$ for the BKT and FQT, and $h = 2$ for FHQT and FQA.

Figure 3 shows an arbitrary BKT, FQT, FHQT, and FQA built on our set of points. Notice that, while in the BKT there is a different pivot per node, in the others there is a different pivot per level, the same for all the nodes of that level.

5.1.1.5. Hybrid. In Shapiro [1977], the use of more than one element per node of the tree is proposed. Those k elements allow eliminating more elements per level at the cost of doing more distance evaluations. The same effect would be obtained if we had a mixture between BKTs and FQTs, so that for k levels we had fixed keys per level, and then we allowed a different key per node of the level $k + 1$, continuing the process recursively on each subtree of the level $k + 1$. The authors conjecture that the pivots should be selected to be outside the clusters.
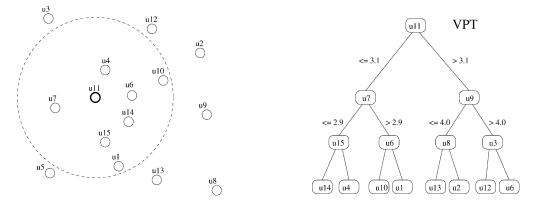
5.1.1.6. Adapting to Continuous Functions. If we have a continuous distance or if it gives too many different values, it is not possible to have a child of the root for any such value. If we did that, the tree would degenerate into a flat tree of height 2, and the search algorithm would be almost like sequential searching for the BKT and

FQT. FHQTs and FQAs do not degenerate in this sense, but they lose their sublinear extra CPU time.

In Baeza-Yates et al. [1994] the authors mention that the structures can be adapted to a continuous distance by assigning a range of distances to each branch of the tree. However, they do not specify how to do this. Some approaches explicitly defined for continuous functions are explained later (VPTs and others), which assign the ranges trying to leave the same number of elements at each class.

5.1.2. Trees for Continuous Distance Functions. We now present the data structures designed for the continuous case. They also can be used for discrete distance functions with virtually no modifications.

5.1.2.1. VPT. The "metric tree" is presented in Uhlmann [1991b] as a tree data structure designed for continuous distance functions. More complete work on the same idea [Yianilos 1993; Chiueh 1994] calls them "vantage-point trees" or VPTs. They build a binary tree recursively, taking any element p as the root and taking the *median* of the set of all distances, $M = median\{d(p, u)/u \in \mathbb{U}\}$. Those

Fig. 4. Example VPT with root u_{11}. We plot the radius M used for the root. For the first levels we show explicitly the radii used in the tree.

elements u such that $d(p, u) \leq M$ are inserted into the left subtree, while those such that $d(p, u) > M$ are inserted into the right subtree. The VPT takes $O(n)$ space and is built in $O(n \log n)$ worst case time, since it is balanced. To solve a query in this tree, we measure $d = d(q, p)$. If $d - r \leq M$ we enter into the left subtree, and if $d + r > M$ we enter into the right subtree (notice that we can enter into both subtrees). We report every element considered that is close enough to the query. See Figure 4.
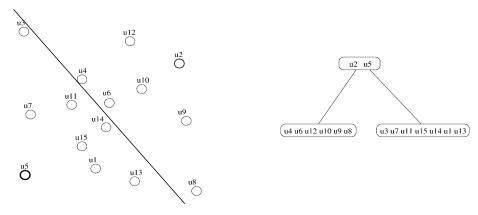
The query complexity is argued to be $O(\log n)$ in Yianilos [1993], but as pointed out, this is true only for very small search radii, too small to be an interesting case.

In trees for discrete distance functions, the exact distance between an element in the leaves and any pivot in the path to the root can be inferred. However, here we only know that the distance is larger or smaller than M. Unlike the discrete case, it is possible that we arrive at an element in a leaf which we do not need to compare, but the tree has not enough information to discover that. Some of those exact distances lost can be stored explicitly, as proposed in Yianilos [1993], to prune more elements before checking them. Finally, Yianilos [1993] considers the problem of pivot selection and argues that it is better to take elements far away from the set.

5.1.2.2. MVPT. The VPT can be extended to m-ary trees by using the $m-1$

uniform percentiles instead of just the median. This is suggested in Brin [1995] and Bozkaya and Ozsoyoglu [1997]. In the latter, the "multi-vantage-point tree" (MVPT) is presented. They propose the use of many elements in a single node, much as in Shapiro [1977]. It can be seen that the space is $O(n)$, since each internal node needs to store the m percentiles but the leaves do not. The construction time is $O(n \log n)$ if we search the m percentiles hierarchically at $O(n \log m)$ instead of $O(mn)$ cost. Bozkaya and Ozsoyoglu [1997] show experimentally that the idea of m-ary trees slightly improves over VPTs (and not in all cases), while a larger improvement is obtained by using many pivots per node. The analysis of query time for VPTs can be extrapolated to MVPTs in a straightforward way.

5.1.2.3. VPF. Another generalization of the VPT is given by the VPF (shorthand for excluded middle vantage point forest) [Yianilos 1999]. This algorithm is designed for radii-limited NN search (an $NN(q)$ query with a maximum radius r^*), but in fact the technique is perfectly compatible with a range search query. The method consists of excluding, at each level, the elements at intermediate distances to their pivot (this is the most populated part of the set): if r_0 and r_n stand for the closest and farthest elements to the pivot p, the elements $u \in \mathbb{U}$ such that $d(p, r_0) + \delta \leq d(p, u) \leq d(p, r_n) - \delta$ are excluded from
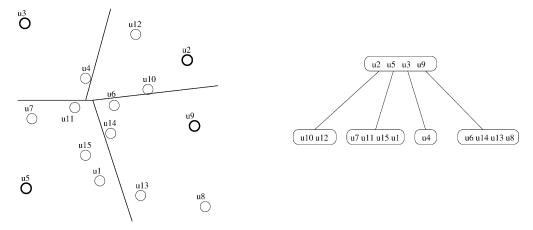
Fig. 5. Example of the first level of a BST or GHT and a query q. Either using covering radii (BST) or hyperplanes (GHT), both subtrees have to be considered in this example.

the tree. A second tree is built with the excluded "middle part" of the first tree, and so on to obtain a forest. With this idea they eliminate the backtracking when searching with a radius $r^* \leq (r_n - r_0 - 2\delta)/2$, and in return they have to search all the trees of the forest. The VPF, of $O(n)$ size, is built using $O(n^{2-\rho})$ time and answers queries in $O(n^{1-\rho} \log n)$ distance evaluations, where $0 < \rho < 1$ depends on r^*. Unfortunately, to achieve $\rho > 0$, r^* has to be quite small.

5.1.2.4. BST. In Kalantari and McDonald [1983], the "bisector trees" (BSTs) are proposed. The BST is a binary tree built recursively as follows. At each node, two "centers" c_1 and c_2 are selected. The elements closer to c_1 than to c_2 go into the left subtree and those closer to c_2 into the right subtree. For each of the two centers, its "covering radius" is stored, that is, the maximum distance from the center to any other element in its subtree. At search time, we enter into each subtree if $d(q, c_i) - r$ is not larger than the covering radius of c_i. That is, we can discard a branch if the query ball (i.e., the hypersphere of radius r centered in the query) does not intersect the ball that contains all the elements inside that branch. In Noltemeier et al. [1992], the "monotonous BST" is proposed, where one of the two elements at each node is indeed the parent center. This makes the covering radii decrease as we move

downward in the tree. Figure 5 illustrates the first step of the tree construction.

5.1.2.5. GHT. Proposed in Uhlmann [1991b], the "generalized-hyperplane tree" (GHT) is identical in construction to a BST. However, the algorithm uses the hyperplane between c_1 and c_2 as the pruning criterion at search time, instead of the covering radius. At search time we enter into the left subtree if $d(q, c_1) - r < d(q, c_2) + r$ and into the right subtree if $d(q, c_2) - r \leq d(q, c_1) + r$. Again, it is possible to enter into both subtrees. In Uhlmann [1991b] it is argued that GHTs could work better than VPTs in high dimensions. The same idea of reusing the parent node is proposed in Bugnion et al. [1993], this time to avoid performing two distance evaluations at each node.

5.1.2.6. GNAT. The GHT is extended in Brin [1995] to an m-ary tree, called GNAT (geometric near-neighbor access tree), keeping the same essential idea. We select, for the first level, m centers $c_1, \ldots, c_m$, and define $\mathbb{U}_i = \{u \in \mathbb{U}, d(c_i, u) < d(c_j, u), \forall j \neq i\}$. That is, $\mathbb{U}_i$ are the elements closer to c_i than to any other c_j. From the root, m children numbered $i = 1 .. m$ are built, each one recursively as a GNAT for $\mathbb{U}_i$. Figure 6 shows a simple example of the first level of a GNAT. Notice the relationship between this idea and a Voronoi-like partition of a vector space [Aurenhammer 1991].

Fig. 6. Example of the first level of a GNAT with $m = 4$.

The search algorithm, however, is quite different. At indexing time, the GNAT stores at each node an $O(m^2)$ size table $range_{ij} = [\min_{u \in \mathbb{U}_j}(c_i, u), \max_{u \in \mathbb{U}_j}(c_i, u)]$, which stores minimum and maximum distances from each center to each class. At search time the query q is compared against some center c_i and then it discards any other center c_j such that $d(q, c_i) \pm r$ does not intersect $range_{i,j}$. All the subtree $\mathbb{U}_j$ can be discarded using the triangle inequality. The process is repeated with random centers until none can be discarded. The search then enters recursively in each nondiscarded subtree. In the process, any center close enough to q is reported.

The authors use a gross analysis to show that the tree takes $O(nm^2)$ space and is built in close to $O(nm \log_m n)$ time. Experimental results show that the GHT is worse than the VPT, which is only beaten with GNATs of arities between 50 and 100. Finally, they mention that the arities of the subtrees could depend on their depth in the tree, but give no clear criteria to do this.

5.1.2.7. VT. The "Voronoi tree" (VT) is proposed in Dehne and Noltemeier [1987] as an improvement over BSTs, where this time the tree has two or three elements (and children) per node. When a new tree node has to be created to hold an inserted element, its closest element from the parent node is also inserted in the new node.

VTs have the property that the covering radius is reduced as we move downwards in the tree, which provides better packing of elements in subtrees. It is shown in Dehne and Noltemeier [1987] that VTs are superior and better balanced than BSTs. In Noltemeier [1989] they show that balanced VTs can be obtained by insertion procedures similar to those of B-trees, a fact later exploited in M-trees (see next).

5.1.2.8. MT. The M-tree (MT) data structure is presented in Ciaccia et al. [1997], aiming at providing dynamic capabilities and good I/O performance in addition to few distance computations. The structure has some resemblances to a GNAT, since it is a tree where a set of representatives is chosen at each node and the elements closer to each representative[2] are organized into a subtree rooted by that representative. The search algorithm, however, is closer to BSTs. Each representative stores its covering radius. At query time, the query is compared to all the representatives of the node and the search algorithm enters recursively into all those that cannot be discarded using the covering radius criterion.
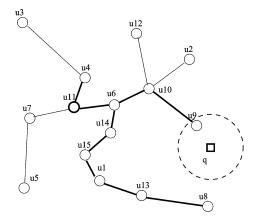
The main difference of the MT is the way in which insertions are handled. An element is inserted into the "best"

---

[2] There are many variants but this is reported as the most effective.

subtree, defined as that causing the subtree covering radius to expand less (zero expansion is the ideal), and in case of ties selecting the closest representative. Finally, the element is added to the leaf node and if the node overflows (i.e., becomes of size $m + 1$) it is split in two and one node element is promoted upwards, as in a B-tree or an R-tree [Guttman 1984]. Hence the MT is a balanced data structure, much as the VP family. There are many criteria to select the representative and to split the node, the best results being obtained by trying a split that minimizes the maximum of the two covering radii obtained. They show experimentally that the MT is resistant to the dimensionality of the space and that it is competitive against R*-trees.

5.1.2.9. SAT. The SAT algorithm (spatial approximation tree) [Navarro 1999] does not use centers to split the set of candidate objects, but rather relies on "spatial" approximation. An element p is selected as the root of a tree, and it is connected to a set of "neighbors" N, defined as a subset of elements $u \in \mathbb{U}$ such that u is closer to p than to any other element in N (note that the definition is self-referential). The other elements (not in $N \cup \{p\}$) are assigned to their closest element in N. Each element in N is recursively the root of a new subtree containing the elements assigned to it.

This allows searching for elements with radius zero by simply moving from the root to its "neighbor" (i.e., connected element) which is closest to the query q. If a radius $r > 0$ is allowed, then we consider that an unknown element $q' \in \mathbb{U}$ is searched with tolerance zero, from which we only know that $d(q, q') \leq r$. Hence, we search as before for q and consider that any distance measure may have an "error" of at most $\pm r$. Therefore, we may have to enter into many branches of the tree (not only the closest one), since the measuring "error" could permit a different neighbor to be the closest one. That is, if $c \in N$ is the closest neighbor of q, we enter into all $c' \in N$ such that $d(q, c') - r \leq d(q, c) + r$. The tree is built in $O(n \log n / \log \log n)$ time, takes
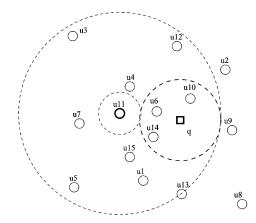

Fig. 7. Example of a SAT and the traversal towards a query q, starting at u_{11}.

$O(n)$ space, and inspects $O(n^{1 - \Theta(1 / \log \log n)})$ elements at query time. Covering radii are also used to increase pruning. Figure 7 shows an example and the search path for a query.

5.1.3. Other Techniques

5.1.3.1. AESA. An algorithm that is close to many of the presented ideas but performs surprisingly better by an order of magnitude is that of Vidal [1986] (called AESA, for approximating eliminating search algorithm). The structure is simply a matrix with the $n(n - 1)/2$ precomputed distances among the elements of $\mathbb{U}$. At search time, they select an element $p \in \mathbb{U}$ at random and measure $r_p = d(p, q)$, eliminating all elements u of $\mathbb{U}$ that do not satisfy $r_p - r \leq d(u, p) \leq r_p + r$. Notice that all the $d(u, p)$ distances are precomputed, so only $d(p, q)$ has been calculated at search time. This process of taking a random pivot among the (not yet eliminated) elements of $\mathbb{U}$ and eliminating more elements from $\mathbb{U}$ is repeated until the candidate set is empty and the query has been satisfied. Figure 8 shows an example with a first pivot u_{11}.

Although this idea seems very similar to FQTs, there are some key differences. The first one, only noticeable in continuous spaces, is that there are no predefined "rings" so that all the intersected rings

Fig. 8. Example of the first iteration of AESA. The points between both rings centered at u_{11} qualify for the next iteration.

qualify (recall Figure 2). Instead, only the minimal necessary area of the rings qualifies. The second difference is that the second element to compare against q is selected from the qualifying set, instead of from the whole set as in FQTs. Finally, the algorithm uses *every* distance computation as a pivoting operation, while FQTs must precompute that decision (i.e., bucket size).

The problem with the algorithm [Vidal 1986] is that it needs $O(n^2)$ space and construction time. This is unacceptably high for all but very small databases. In this sense the approach is close to Sasha and Wang [1990], although in this latter case they may take fewer distances and bound the unknown ones. AESA is experimentally shown to have $O(1)$ query time.

5.1.3.2. LAESA and Variants. In a newer version of AESA, called LAESA (for linear AESA), Micó et al. [1994] propose using k fixed pivots, so that the space and construction time is $O(kn)$. In this case, the only difference with an FHQT is that fixed rings are not used, but the exact set of elements in the range is retrieved. FHQT uses fixed rings to reduce the extra CPU time, while in this case no such algorithm is given. In LAESA, the elements are simply linearly traversed, and those that cannot be eliminated after considering the k pivots are directly compared against the query.

A way to reduce the extra CPU time is presented later in Micó et al. [1994], which builds a GHT-like structure using the same pivots. The algorithm is argued to be sublinear in CPU time. Alternative search structures to reduce CPU time not losing information on distances are presented in Nene and Nayar [1997] and Chávez et al. [1999], where the distances to each pivot are sorted separately so that the relevant range $[d(q, p)-r, d(q, p)+r]$ can be binary searched.[3] Extra pointers are added to be able to trace an element across the different orderings for each pivot (this needs more space, however).

5.1.3.3. Clustering and LCs. Clustering is a very wide area with many applications [Jain and Dubes 1988]. The general goal is to divide a set into subsets of elements close to each other in the same subset. A few approaches to index metric spaces based on clustering exist.

One of them, LC (for "list of clusters") [Chávez and Navarro 2000] chooses an element p and finds the m closest elements in U. This is the cluster of p. The process continues recursively with the remaining elements until a list of $n/(m + 1)$ clusters is obtained. The covering radius $cr()$ of each cluster is stored. At search time the clusters are inspected one by one. If $d(q, p_i) - r > cr(p_i)$ we do not enter the cluster of p_i, otherwise we verify it exhaustively. If $d(q, p_i) + r < cr(p_i)$ we do not need to consider subsequent clusters. The authors report unbeaten performance on high dimensional spaces, at the cost of $O(n^2/m)$ construction time.

5.2. Nearest-Neighbor Queries

We have concentrated in range search queries up to now. This is because, as we show in this section, most of the existing solutions for NN-queries are built systematically over range searching techniques, and indeed can be adapted to any of the data structures presented (despite having been originally designed for specific ones).

---

[3] Although Nene and Nayar [1997] consider only vector spaces, the same technique can be used here.

5.2.1. Increasing Radius. The simplest NN search algorithm is based on using a range searching algorithm as follows. Search q with fixed radii $r = a^i \varepsilon$ $(a > 1)$, starting with $i = 0$ and increasing it until at least the desired number of elements (1 or k) lies inside the search radius $r = a^i \varepsilon$.

Since the complexity of the range searching normally grows sharply on the search radius, the cost of this method can be very close to the cost of a range searching with the appropriate r (which is not known in advance). The increasing steps can be made smaller $(a \to 1)$ to avoid searching with a radius much larger than necessary.

5.2.2. Backtracking with Decreasing Radius. A more elaborate technique is as follows. We first explain the search for the closest neighbor. Start the search on any data structure using $r^* = \infty$. Each time q is compared to some element p, update the search radius as $r^* \leftarrow \min(r^*, d(q, p))$ and continue the search with this reduced radius. This has been proposed, for example, for BKTs and FQTs [Burkhard and Keller 1973; Baeza-Yates et al. 1994] as well as for vector spaces [Roussopoulos et al. 1995].

As closer and closer elements to q are found, we search with smaller radius and the search becomes cheaper. For this reason it is important to try to find quickly elements that are close to the query (which is unimportant in range queries). The way to achieve this is dependent on the particular data structure. For example, in BKTs and FQTs we can begin at the root and measure $i = d(p, q)$. Now, we consider the edges labeled $i, i-1, i+1, i-2, i+2$, and so on, and proceed recursively in the children (other heuristics may work better). Therefore, the exploration ends just after considering the branch $i + r^*$ (r^* is reduced along the process). At the end r^* is the distance to the closest neighbors and we have already seen all of them.

$NN_k(q)$ queries are solved as an extension of the above technique, where we keep the k elements seen that are closest to q and set r^* as the maximum distance between those elements and q (clearly we are not interested in elements farther away than the current kth closest element). Each time a new element is seen whose distance is relevant, it is inserted as one of the k nearest neighbors known up to now (possibly displacing one of the old candidates from the list) and r^* is updated. In the beginning we start with $r^* = \infty$ and keep this value until the first k elements are found.

A variant of this type of query is the *limited radius* NN searching. Here we start with the maximum expected distance between the query element and its nearest neighbor. This type of query has been the focus of Yianilos [1999, 2000].

5.2.3. Priority Backtracking. The previous technique can be improved by a smarter selection of which elements to consider first. For clarity we consider backtracking in a tree, although the idea is general. Instead of following the normal backtracking order of the range query, modifying at most the order in which the subtrees are traversed, we give much more freedom to the traversal order. The goal is to increase the probability of quickly finding elements close to q and therefore reduce r^* fast. This technique has been used in vector and metric spaces several times [Uhlmann 1991a; Hjaltason and Samet 1995; Ciaccia et al. 1997].

At each point of the search we have a set of possible subtrees that can be traversed (not necessarily all at the same level). For each subtree, we know a lower bound to the distance between q and any element of the subtree. We first select subtrees with the smallest lower bound. Once a subtree has been selected we compare q against its root, update r^* and the candidates for output if necessary, and determine which of the children of the considered root deserve traversal. Unlike the normal backtracking, those children are not immediately traversed but added to the set of subtrees that have to be traversed at some moment. Then we select again a subtree from the set.

The best way to implement this search is with a priority queue ordered by the heuristic "goodness," where the subtrees

are inserted and removed. We start with an empty queue where we insert the root of the tree. Then we repeat the step of removing the most promising subtree, processing it, and inserting the relevant subtrees until the lower bound of the best subtree exceeds r^*.

If applied to a BKT or a FQT, this method yields the same result as the previous section, but this technique is superior for dealing with continuous distances.

5.2.4. Specific NN Algorithms. The techniques described above cover almost all the existing proposals for solving NN-queries. The only exception we are aware of was presented in Clarkson [1999], which is a GNAT-like data structure where the points are inserted into more than one subtree to limit backtracking (hence the space requirement is superlinear).

After selecting the representatives for the root of the tree, each element u is not only inserted into the subtree of its closest representative p, but also in the tree of any other representative p' such that $d(u, p') \leq 3d(u, p)$. At search time, the query q enters not only into its nearest representative p but also into every other representative p' such that $d(q, p') \leq 3d(q, p)$. As shown in Clarkson [1999] this is enough to guarantee that the nearest neighbor will be reached.

By using subsets of size $n^{1/2^{k+1}}$ at depth k in the tree, the search time is polylogarithmic in n and the space requirement is $O(n \text{ polylog } n)$ if some conditions hold in the metric space.

5.3. Extensions

We cover in this section the work that has been pursued on extensions of the basic problems or in alternative models. None of these are the main focus of our survey.

5.3.1. Dynamic Capabilities. Many of the data structures for metric spaces are designed to be built on a static data set. In many applications this is not reasonable because elements have to be inserted and deleted dynamically. Some data structures tolerate insertions well, but not deletions.

We first consider insertions. Among the structures that we have surveyed, the least dynamic is SAT, which needs full knowledge of the complete set at index construction time and has difficulty in handling later insertions (some solutions are described in Navarro and Reyes [2001]). The VP family (VPT, MVPT, VPF) has the problem of relying on global statistics (such as the median) to build the tree, so later insertions can be performed but the performance of the structure may deteriorate. Finally, the FQA needs, in principle, insertion in the middle of an array, but this can be handled by using standard techniques. All the other data structures can handle insertions in a reasonable way. There are some structure parameters that may depend on n and thus require periodical structural reorganization, but we disregard this issue here (e.g., adding or removing pivots is generally problematic).

Deletion is a little more complicated. In addition to the above structures, which present the same problems as for insertion, BKTs, GHTs, BSTs, VTs, GNATs, and the VP family cannot tolerate deletion of an internal tree node because it plays an essential role in organizing the subtree. Of course this can be handled as just marking the node as removed and actually keeping it for routing purposes, but the quality of the data structure is affected over time.

Therefore, the only structures that fully support insertions and deletions are in the FQ family (FQT, FQHT, FQA, since there are no truly internal nodes), AESA and LAESA approaches (since they are just vectors of coordinates), the MT (which is designed with dynamic capabilities in mind and whose insertion/deletion algorithms are reminiscent of those of the B-tree), and a variant of GHTs designed to support dynamic operations [Verbarg 1995]. The analysis of this latter structure shows that dynamic insertions can be done in $O(\log^2 n)$ amortized worst case time, and that deletions can be done at similar cost under some restrictions.

5.3.2. I/O Considerations. Most of the research on metric spaces deals with reducing the number of distance evaluations or at most the total CPU time. However, depending on the application, the I/O cost may play an important role. As most of the research on metric spaces has focused on algorithms to discard elements, I/O considerations have been normally left aside.

The only exception to the rule is MT, designed specifically for secondary memory. The tree nodes in the MT are to be stored in a single disk page (indeed, the MT does not fix an arity but rather a node capacity in bytes). Earlier balanced trees exist (such as the VP family), but the purpose of this balancing is to keep the extra CPU costs low. As we show later, unbalanced data structures perform much better in high dimensions and it is unlikely that the reduced CPU costs may play an important role. The purpose of balancing the MT, on the other hand, is to keep I/O costs low, and depending on the application this may be even more important than the number of distance evaluations.

Other data structures could probably be adapted to perform well in secondary memory, but the authors simply have not considered the problem. For instance, it is not hard to imagine strategies for the tree data structures to pack "compact" subtrees in disk pages, so as to make as good use as possible of a page that is read from disk. When a subtree grows larger than a page it is split into two pages of approximately the same number of nodes. Of course, a B-tree-like scheme such as that of MT has to be superior in this respect. Finally, array-oriented approaches such as FQA, AESA, and LAESA are likely to read all the disk pages of the index for each query, and hence have bad I/O performance.

5.3.3. Approximate and Probabilistic Algorithms. For the sake of a complete overview we include a brief description of an important branch of similarity searching, where a relaxation on the query precision is allowed to obtain a speedup in the query time complexity. This is reason-

able in some applications because the metric space modelization already involves an approximation to the true answer (recall Section 2), and therefore a second approximation at search time may be acceptable.

Additionally to the query one specifies a precision parameter ε to control how far away (in some sense) we want the outcome of the query from the correct result. A reasonable behavior for this type of algorithm is to asymptotically approach the correct answer as ε goes to zero, and complementarily to speed up the algorithm, losing precision, as ε moves in the opposite direction.

This alternative to *exact similarity searching* is called *approximate similarity searching*, and encompasses approximate and probabilistic algorithms. We do not cover them in depth here but present a few examples. Approximation algorithms for similarity searching are considered in depth in White and Jain [1996].

As a first example, we mention an approximate algorithm for NN search in real vector spaces using any L_s metric by Arya et al. [1994]. They propose a data structure, the BBD-tree, inspired in kd-trees, that can be used to find "$(1 + \varepsilon)$ nearest neighbors": instead of finding u such that $d(u,q) \leq d(v,q) \; \forall v \in \mathbb{U}$, they find an element u^*, an $(1+\varepsilon)$-nearest neighbor, differing from u by a factor of $(1 + \varepsilon)$, that is, u^* such that $d(u^*,q) \leq (1+\varepsilon)d(v,q) \; \forall v \in \mathbb{U}$.

The essential idea behind this algorithm is to locate the query q in a cell (each leaf in the tree is associated with a cell in the decomposition). Every point inside the cell is processed to obtain the current nearest neighbor (u). The search stops when no promising cells are encountered, that is, when the radius of any ball centered at q and intersecting a nonempty cell exceeds the radius $d(q,p)/(1+\varepsilon)$. The query time is $O(\lceil 1 + 6k/\varepsilon \rceil^k k \log n)$.

A second example is a probabilistic algorithm for vector spaces [Yianilos 2000]. The data structure is like a standard kd-tree, using "aggressive pruning" to improve the performance. The idea is to increase the number of branches pruned
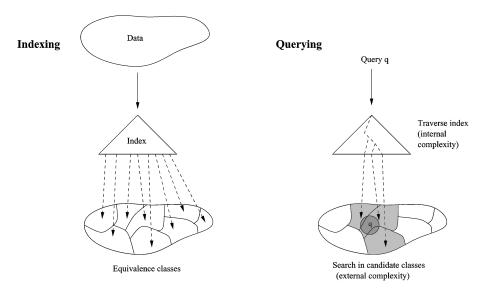
Fig. 9. The unified model for indexing and querying metric spaces.

at the expense of losing some candidate points in the process. This is done in a controlled way, so the probability of success is always known. In spite of the vector space focus of the algorithm, it could be generalized to metric spaces as well. The data structure is useful for finding only limited radius nearest neighbors, that is, neighbors within a fixed distance to the query.

Finally, an example of a probabilistic NN algorithm for general metric spaces is that of Clarkson [1999]. The original intention is to build a Voronoi-like data structure on a metric space. As this is not possible in general because there is not enough knowledge of the characteristics of the queries that will come later [Navarro 1999], Clarkson [1999] chooses to have a "training set" of queries and to build a data structure able to answer correctly only queries belonging to the training set. The idea is that this is enough to answer correctly an arbitrary query with high probability. Under some probabilistic assumptions on the distribution of the queries, it is shown that the probability of not finding the nearest neighbor is $O(\log n)^2/K)$, where K can be arbitrarily large at the expense of $O(Kn\rho)$ space and $O(K\rho \log n)$ expected search

time. Here ρ is the logarithm of the ratio between the farthest and the nearest pairs of points in the union of $\mathbb{U}$ and the training set.

6. A UNIFYING MODEL

At first sight, the indexes and the search algorithms seem to emerge from a great diversity, and different approaches are analyzed separately, often under different assumptions. Currently, the only realistic way to compare two different algorithms is to apply them to the same data set.

In this section we make a formal introduction to our unifying model. Our intention is to provide a common framework to analyze all the existing approaches to proximity searching. As a result, we are able to capture the similarities of apparently different approaches. We also obtain truly new ways of viewing the problem.

The conclusion of this section can be summarized in Figure 9. All the indexing algorithms partition the set $\mathbb{U}$ into subsets. An index is built that allows determining a set of candidate subsets where the elements relevant to the query can appear. At query time, the index is searched to find the relevant subsets (the cost of doing this is called "internal complexity") and those

subsets are checked exhaustively (which corresponds to the "external complexity" of the search).

The last two subsections describe the two main approaches to similarity searching in abstract terms.

6.1. Equivalence Relations

The relevance of equivalence classes for this survey comes from the possibility of partitioning a metric space so that a new metric space is derived from the quotient set. Readers familiar with equivalence relations can safely skip this short section.

Given a set $\mathbb{X}$, a *partition* $\pi(\mathbb{X}) = \{\pi_1, \pi_2, \dots\}$ is a collection of pairwise disjoint subsets whose union is $\mathbb{X}$; that is, $\cup \pi_i = \mathbb{X}$ and $\forall i \neq j, \pi_i \cap \pi_j = \emptyset$.

A relation, denoted by $\sim$, is a subset of the crossproduct $\mathbb{X} \times \mathbb{X}$ (the set of ordered pairs) of $\mathbb{X}$. Two elements x, y are said to be related, denoted by $x \sim y$, if the pair (x, y) is in the subset. A relation $\sim$ is said to be an *equivalence relation* if it satisfies, for all $x, y, z \in \mathbb{X}$, the properties of reflexivity $(x \sim x)$, symmetry $(x \sim y \Leftrightarrow y \sim x)$, and transitivity $(x \sim y \land y \sim z \Rightarrow x \sim z)$.

Every partition $\pi(\mathbb{X})$ induces an equivalence relation $\sim$ and, conversely, every equivalence relation induces a partition: two elements are related if they belong to the same partition element. Every element π_i of the partition is then called an *equivalence class*. An equivalence class is often named after one of its representatives (any element of π_i can be taken as a representative). An alternative definition of an equivalence class of an element x is the set of all y such that $x \sim y$. We denote the equivalence class of x as $[x] = \{y, x \sim y\}$.

Given the set $\mathbb{X}$ and an equivalence relation $\sim$, we obtain the quotient $\pi(\mathbb{X}) = \mathbb{X}/\sim$. It indicates the set of equivalence classes (or just classes), obtained when applying the equivalence relation to the set $\mathbb{X}$.

6.2. Indexing and Partitions

The equivalence classes in the quotient set $\pi(\mathbb{X})$ of a metric space $\mathbb{X}$ can be considered themselves as objects in a new metric space, provided we define a distance function in $\pi(\mathbb{X})$.

We introduce a new function $D_0 : \pi(\mathbb{X}) \times \pi(\mathbb{X}) \to \mathbb{R}$ now defined in the quotient.

Definition 1. Given a metric space $(\mathbb{X}, d)$ and a partition $\pi(\mathbb{X})$, the *extension* of d to $\pi(\mathbb{X})$ is defined as $D_0([x], [y]) = \inf_{x \in [x], y \in [y]} \{d(x, y)\}$.

D_0 gives the maximum possible values that keep the mapping *contractive* (i.e., $D_0([x], [y]) \leq d(x, y)$ for any x, y). Unfortunately, D_0 does not satisfy the triangle inequality, just $(p1)$ to $(p3)$, and in most cases $(p4)$ (recall Section 3.1). Hence, D_0 itself is not suitable for indexing purposes.
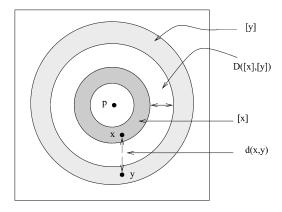
However, we can use any metric D that lower bounds D_0 (i.e., $D([x], [y]) \leq D_0([x], [y])$). Since D is a metric, $(\pi(\mathbb{X}), D)$ is a metric space and therefore we can make queries in $\pi(\mathbb{X})$ in the same way we have done in $\mathbb{X}$. We redefine the outcome of a query in $\pi(\mathbb{X})$ as $([q], r)_D = \{u \in \mathbb{U}, D([u], [q]) \leq r\}$ (although formally we should retrieve classes, not elements).

Since the mapping is contractive $(D([x], [y]) \leq d(x, y))$ we can convert one search problem into another, it is hoped simpler, search problem. For a given query $(q, r)_d$ we find out to which equivalence class the query q belongs (i.e., $[q]$). Then, using the new distance function D the query is transformed into $([q], r)_D$. As the mapping is contractive, we have $(q, r)_d \subseteq ([q], r)_D$. That is, $([q], r)_D$ is indeed a candidate list, so it is enough to perform an exhaustive search on that candidate list (now using the original distance d), to obtain the actual outcome of the query $(q, r)_d$.

Our main thesis is that the above procedure is in fact used in virtually every indexing algorithm (recall Figure 9).

PROPOSITION. *All the existing indexing algorithms for proximity searching consist of building an equivalence relation, so that at search time some classes are discarded and the others are exhaustively searched.*

As we show shortly, the most important tradeoff when designing the partition is to

Fig. 10. Two points x and y, and their equivalence classes (the shaded rings). D gives the minimal distance among rings, which lower bounds the distance between x and y.

balance the cost to find $([q], r)_D$ and the cost to verify this candidate list.

In Figure 10 we can see a schematic example of the idea. We divide the space in several regions (equivalence classes). The objects inside each region become indistinguishable. We can consider them as elements in a new metric space. To find the answer, instead of exhaustively examining the entire dictionary we just examine the classes that contain potentially interesting objects. In other words, if a class can contain an element that should be returned in the outcome of the query, then the class will be examined (see also the rings considered in Figure 2).

We recall that this property is not enough for an arbitrary NN search algorithm to work (since the mapping would have to preserve proximity instead), but most existing algorithms for NN are based on range queries (recall Section 5.2), and these algorithms can be applied as well.

Some examples may help to understand the above definitions, for both the concept of equivalence relation and the obtained distance function.

Example 1. Say that we have an arbitrary reference pivot $p \in \mathbb{X}$ and the equivalence relation is given by $x \sim y \Leftrightarrow d(p, x) = d(p, y)$. In this case $D([x],[y]) = |d(x, p) - d(y, p)|$ is a safe lower bound for the D_0 distance (guaranteed by the triangle inequality). For a query of the form

$(q, r)_d$ the candidate list $([q], r)_D$ consists of all elements x such that $D([q],[x]) \le r$, or which is the same, $|d(q, p) - d(x, p)| \le r$. Graphically, this distance represents a *ring* centered at p containing a ball centered at q with radius r (recall Figures 10 and 8). This is the familiar rule used in many independent algorithms to trim the space.

Example 2. As explained, the similarity search problem was first introduced in vector spaces, and the very first family of algorithms used there was based on a partition operation. These algorithms were called *bucketing* methods, and consisted of the construction of cells or buckets [Bentley et al. 1980] Searching for an arbitrary point in $\mathbb{R}^k$ is converted into an exhaustive search in a finite set of cells. The procedure uses two steps: they find which cell the query point belongs to and build a set of candidate cells using the query range; then they inspect this set of candidate cells exhaustively to find the actual points inside the query range.[4] In this case the equivalence classes are the cells, and the tradeoff is that the larger the cells, the cheaper it is to find the appropriate ones, but the more costly is the final exhaustive search.

6.3. Coarsening and Refining a Partition

We start by defining the concepts of refinement and coarsening.

Definition 2. Let $\sim_1$ and $\sim_2$ be two equivalence relations over a set $\mathbb{X}$. We say that $\sim_1$ is a refinement of $\sim_2$ or that $\sim_2$ is a coarsening of $\sim_1$ if for any pair $x, y \in \mathbb{X}$ such that $x \sim_1 y$ it holds $x \sim_2 y$. The same terms can be applied to the corresponding partitions $\pi^1(\mathbb{X})$ and $\pi^2(\mathbb{X})$.

Refinement and coarsening are important concepts for the topic we are discussing. The following lemma shows the effect of coarsening on the effectiveness of the partition for searching purposes.

---

[4]The algorithm is in fact a little more sophisticated because they try to find the nearest neighbor of a point. However, the version presented here for range queries is in the same spirit as the original one.

LEMMA 1. *If $\sim_1$ is a coarsening of $\sim_2$ then their extended distances D_1 and D_2 have the property $D_1([x],[y]) \leq D_2([x],[y])$.*

PROOF. Let us denote $[x]_i$ and $[y]_i$ the equivalence classes of x and y under equivalence relation $\sim_i$. Then

$$
\begin{aligned}
D_1([x],[y]) &= \inf_{x\in[x]_1, y\in[y]_1}\{d(x,y)\} \\
&\leq \inf_{x\in[x]_2, y\in[y]_2}\{d(x,y)\} = D_2([x],[y]),
\end{aligned}
$$

since $[x]_2 \subseteq [x]_1$ and $[y]_2 \subseteq [y]_1$. □

An interesting idea arising from the above lemma is to build a hierarchy of coarsening operations. Using this hierarchy we could proceed downwards from a very coarse level building a candidate list of equivalence classes of the next level. This candidate list will be refined using the next distance function and so on until we reach the bottom level.

6.4. Discriminative Power

As sketched previously, most indexing algorithms rely on building an equivalence relation. The corresponding search algorithms have the following parts.

1. Find the classes that may be relevant for the query.
2. Exhaustively search all the elements of these classes.

The first part involves performing some evaluations of the d distance, as shown in Example 1 above. It may also involve some extra CPU time (which although not the central point in this article, must be kept reasonable). The second part consists of directly comparing the query against the candidate list. The following definition gives a name to both parts of the search cost.

Definition 3. Let A be a search algorithm over $(\mathbb{X}, d)$ based on a mapping to $(\pi(\mathbb{X}), D)$, and let $(q,r)_d$ be a range query. Then the internal complexity of A is the number of evaluations of d necessary to compute $([q], r)_D$, and the external complexity is $|([q], r)_D|$.

We recall that $|([q], r)_D|$ refers to the number of elements in the original metric space, not the number of classes retrieved.

There is a concept related to the external complexity of a search algorithm, which we define next.

Definition 4. The discriminative power of a search algorithm based on a mapping from $(\mathbb{X}, d)$ to $(\pi(\mathbb{X}), D)$, with regard to a query $(q,r)_d$ of nonempty outcome, is defined as $|(q,r)_d|/|([q], r)_D|$.

Although the definition depends on q and r, we can speak in general terms of the discriminative power by averaging over the qs and rs of interest. The discriminative power serves as an indicator of the performance or fitness of the equivalence relation.

In general, it will be more costly to have more discriminative power. The indexing scheme needs to find a balance between the complexity to find the relevant classes and the discriminative power.

Let us consider Example 1. The internal complexity is one distance evaluation (the distance from q to p), and the external complexity will correspond to the number of elements that lie in the selected ring. We could intersect it with more rings (increasing internal complexity) to reduce the external complexity.

The tradeoff is partially formalized with this lemma.

LEMMA 2. *If A_1 and A_2 are search algorithms based on equivalence relations $\sim_1$ and $\sim_2$, respectively, and $\sim_1$ is a coarsening of $\sim_2$, then A_1 has higher external complexity than A_2.*

PROOF. We have to show that, for any r, $([q], r)_{D_2} \subseteq ([q], r)_{D_1}$. But this is clear, since $D_1([x],[y]) \leq D_2([x],[y])$ implies $([q], r)_{D_2} = \{y \in \mathbb{U}, \ D_2([q],[y]) \leq r\} \subseteq \{y \in \mathbb{U}, \ D_1([q],[y]) \leq r\} = ([q], r)_{D_1}$. □

Although having more discriminative power normally costs more internal evaluations, one can make better or worse use of the internal complexity. We elaborate more on this next.
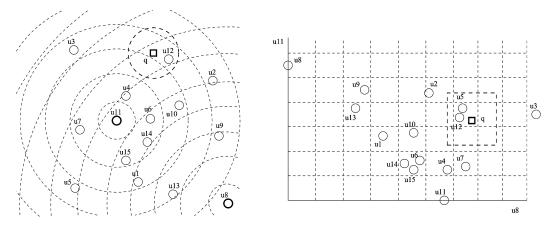
Fig. 11. An equivalence relation induced by intersecting rings centered in two pivots, and how a query is transformed.

6.5. Locality of a Partition

The equivalence classes can be thought of as a set of nonintersecting cells in the space, where every element inside a given cell belongs to the same equivalence class. However, the mathematical definition of an equivalence class is not confined to a single cell. We define "locality" as a property of a partition that stands for how much the classes resemble cells.

Definition 5. The *nonlocality* of a partition $\pi(\mathbb{X}) = \{\pi_1, \pi_2, \ldots\}$ with respect to a finite dictionary $\mathbb{U}$ is defined as $\max_i\{\max_{x,y \in \pi_i \cap \mathbb{U}} d(x, y)\}$, that is, as the maximum distance between elements lying in the same class.

We say that a partition is "local" or "nonlocal" meaning that it has high or low locality. Figure 11 shows an example of a nonlocal partition (u_5 and u_{12} lie in separate fragments of a single class). It is natural to expect more discriminative power from a local partition than from a nonlocal one. This is because in a nonlocal partition the candidate list tends to contain elements actually far away from the query. Notice that in Figure 11 the locality would improve sharply if we added a third pivot. In a vector space of k dimensions, it suffices to consider $k + 1$ pivots in general position[5] to obtain a highly local partition.

In general metric spaces we can also take a sufficient number of pivots so as to obtain highly local partitions.

However, obtaining local partitions may be expensive in internal complexity and not enough to achieve low external complexity, otherwise the bucketing method for vector spaces [Bentley et al. 1980] explained in Example 2 would have excellent performance. Even with such a local partition and assuming uniformly distributed data, a number of empty cells are verified, whose volume grows exponentially with the dimension. We return later to this issue.

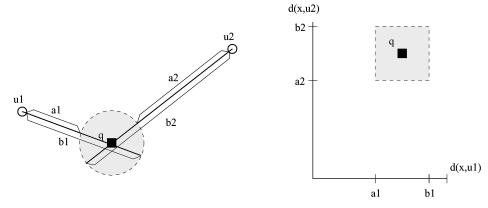## 6.6. The Pivot Equivalence Relation

A large class of algorithms to build the equivalence relations is based on *pivoting*. This consists of considering the distances between an element and a number of preselected "pivots" (i.e., elements of $\mathbb{U}$ or even $\mathbb{X}$, also called reference points, vantage points, keys, queries, etc. in the literature).

The equivalence relation is defined in terms of the distances of the elements to the pivots, so that two elements are equivalent if they are at the same distance from all the pivots. If we consider one pivot p, then this equivalence relation is

$$x \sim_p y \Leftrightarrow d(x, p) = d(y, p).$$

---

[5] That is, not lying on a $(k - 1)$-hyperplane.

Fig. 12.   Mapping from a metric space onto a vector space under the L_∞ metric, using two pivots.

The equivalence classes correspond to the intuitive notion of the family of sphere shells with center p. Points falling in the same sphere shell (i.e., at the same distance from p) are equivalent from the viewpoint of p.

The above equivalence relation is easily generalized to k pivots.

Definition 6.   The *pivot equivalence relation* based on elements $\{p_1, \ldots, p_k\}$ (the k pivots) is defined as

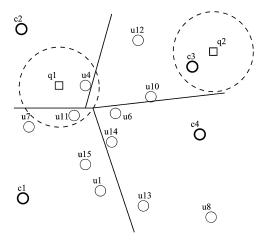$$x \sim_{\{p_i\}} y \Leftrightarrow \forall i, \ d(x, p_i) = d(y, p_i).$$

A graphical representation of the class in the general case corresponds to the intersection of several sphere shells centered at the points p_i (recall Figure 11).

The distance $d(x, y)$ cannot be smaller than $|d(x, p) - d(y, p)|$ for any element p, because of the triangle inequality. Hence $D([x], [y]) = |d(x, p) - d(y, p)|$ is a safe lower bound to the D_0 function corresponding to the class of sphere shells centered in p. With k pivots, this becomes $D([x], [y]) = \max_i\{|d(x, p_i) - d(y, p_i)|\}$. This D distance lower bounds d and hence can be used as our distance in the quotient space.

Alternatively, we can consider the equivalence relation as a projection to the vector space $\mathbb{R}^k$, k being the number of pivots used. The ith coordinate of an element is the distance of the element to the ith pivot. Once this is done, we can identify points in $\mathbb{R}^k$ with elements in the original space with the L_∞ distance. As we have described in Section 6, the indexing algorithm will consist of finding the set of equivalence classes such that they fall inside the radius of the search when using D in the quotient space. In this particular case for a query of the form $(q, r)_d$ we have to find the candidate list as the set $([q], r)_D$, that is, the set of equivalence classes $[y]$ such that $D([q], [y]) \leq r$. In other words, we want the set of objects y such that $\max_i\{|d(q, p_i) - d(y, p_i)|\} \leq r$. This is equivalent to search with the L_∞ distance in the vector space $\mathbb{R}^k$ where the equivalence classes are projected. Figure 12 illustrates this concept (Figure 11 is also useful).

Yet a third way to see the technique, less formal but perhaps more intuitive, is as follows. To check if an element $u \in \mathbb{U}$ belongs to the query outcome, we try a number of random pivots p_i. If, for any such p_i, we have $|d(q, p_i) - d(u, p_i)| > r$, then by the triangle inequality we know that $d(q, u) > r$ without the need to actually evaluate $d(q, u)$. At indexing time we precompute the $d(u, p_i)$ values and at search time we compute the $d(q, p_i)$ values. Only those elements u that cannot be discarded by looking at the pivots are actually checked against q.

Fig. 13. A compact partition using four centers and two query balls intersecting some classes. Note that the class of c_3 could be excluded from consideration for q_1 by using the covering radius criterion but not the hyperplane criterion, whereas the opposite happens to discard c_4 for q_2.

6.7. The Compact Equivalence Relation

A different type of equivalence relation, used by another large class of algorithms, is defined with respect to the proximity to a set of elements (which we call "centers" to distinguish them from the pivots of the previous section).

Definition 7. The *compact equivalence relation* based on $\{c_1, \ldots, c_m\}$ (the centers) is

$$x \sim_{\{c_i\}} y \Leftrightarrow closest(x, \{c_i\}) = closest(y, \{c_i\}),$$

where $closest(z, S) = \{w \in S, \forall w' \in S, d(z, w) \leq d(z, w')\}$. The associated partition is called a compact partition.

That is, we divide the space with one partition element for each c_i and the class of c_i is that of the points that have c_i as their closest center. Figure 13 shows an example in $(\mathbb{R}^2, L_2)$. In particular, note that we can select $\mathbb{U}$ as the set of centers, in which case the partition has optimal locality. Even if $\{c_i\}$ is not $\mathbb{U}$, compact partitions have good locality.

In vector spaces the compact partition coincides with the Voronoi partition if the whole set $\mathbb{U}$ is the set of centers. Its associated concept, the "Delaunay tesselation," is a graph whose nodes are the elements of $\mathbb{U}$ and whose edges connect nodes whose classes share a border. The Delaunay tesselation is the basis of very good algorithms for proximity searching [Aurenhammer 1991; Yao 1990]. For example, an $O(\log n)$ time NN algorithm exists in two dimensions. Unfortunately, this algorithm does not generalize efficiently to more than two dimensions. The Delaunay tesselation, which has $O(n)$ edges in two dimensions, can have $O(n^2)$ edges in three and more dimensions.

In a general metric space, the $D_0([x], [y])$ distance in the space $\pi(\mathbb{X})$ of the compact classes is, as before, the smallest distance between points $x \in [x]$ and $y \in [y]$. To find $[q]$ we basically need to find the nearest neighbor of q in the set of centers. The outcome of the query $([q], r)_{D_0}$ is the set of classes intersected by a query ball (see Figure 13).
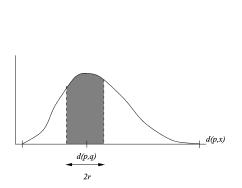
A problem in general metric spaces is that it is not easy to bound a class so as to determine whether the query ball intersects it. From the many possible criteria, two are the most popular.
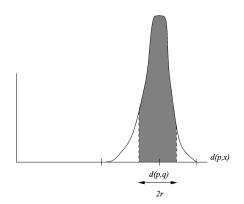
6.7.1. Hyperplane Criterion. This is the most basic one and the one that best expresses the idea of compactness. In essence, if c is the center of the class $[q]$ (i.e., the center closest to q), then (1) the query ball of course intersects $[c]$; (2) the query ball does not intersect $[c_i]$ if $d(q, c) + r < d(q, c_i) - r$. Graphically, if the query ball does not intersect the hyperplane dividing its closest neighbor and another center c_i, then the ball is totally outside the class of c_i.

6.7.2. Covering Radius Criterion. This tries to bound the class $[c_i]$ by considering a ball centered at c_i that contains all the elements of $\mathbb{U}$ that lie in the class.

Definition 8. The *covering radius* of c for $\mathbb{U}$ is $cr(c) = \max_{u \in [c] \cap \mathbb{U}} d(c, u)$.

Now it is clear that we can discard c_i if $d(q, c_i) - r > cr(c_i)$.

Fig. 14. A low-dimensional (left) and high-dimensional (right) histogram of distances, showing that on high dimensions virtually all the elements become candidates for exhaustive evaluation. Moreover, we should use a larger *r* in the second plot in order to retrieve some elements.

7. THE CURSE OF DIMENSIONALITY

As explained, one of the major obstacles for the design of efficient search techniques on metric spaces is the existence and ubiquity in real applications of the so-called *high-dimensional spaces*. Traditional indexing techniques for vector spaces (such as *kd*-trees) have an exponential dependency on the representational dimension of the space (as the volume of a box or hypercube containing the answers grows exponentially with the dimension).

More recent indexing techniques for vector spaces and those for generic metric spaces can get rid of the representational dimension of the space. This makes a big difference in many applications that handle vector spaces of high representational dimension but low intrinsic dimension (e.g., a plane immersed in a 50-dimensional vector space, or simply clustered data). However, in some cases even the intrinsic dimension is very high and the problem becomes intractable for exact algorithms, and we have to resort to approximate or probabilistic algorithms (Section 5.3.3).

Our aim in this section is (*a*) to show that the concept of intrinsic dimensionality can be conceived even in a general metric space; (*b*) to give a quantitative definition of the intrinsic dimensionality; (*c*) to show analytically the reason for the so-called "curse of dimensionality;" and (*d*) to discuss the effects of pivot selection tech-

niques. This is partially based on previous work [Chávez and Navarro 2001a].

7.1. Intrinsic Dimensionality

Let us start with a well-known example. Consider a distance such that $d(x, x) = 0$ and $d(x, y) = 1$ for all $x \neq y$. Under this distance (in fact an equality test), we do not obtain any information from a comparison except that the element considered is or is not our query. It is clear that it is not possible to avoid a sequential search in this case, no matter how smart our indexing technique is.

Let us consider the *histogram* of distances between points in the metric space $\mathbb{X}$. This can be approximated by using the dictionary $\mathbb{U}$ as a random sample of $\mathbb{X}$. This histogram is mentioned in many papers, for example, Brin [1995], Chávez and Marroquín [1997], and Ciaccia et al. [1998a], as a fundamental measure related to the intrinsic dimensionality of the metric space. The idea is that, as the space has higher intrinsic dimension, the mean μ of the histogram grows and its variance σ^2 is reduced. Our previous example is an extreme case.

Figure 14 gives an intuitive explanation of why the search problem is harder when the histogram is concentrated. If we consider a random query q and an indexing scheme based on random pivots, then the possible distances between q and a pivot p are distributed according to the histogram

of the figure. The elimination rule says that we can discard any point u such that $d(p, u) \notin [d(p, q) - r, d(p, q) + r]$. The grayed areas in the figure show the points that we *cannot* discard. As the histogram is more and more concentrated around its mean, fewer points can be discarded using the information given by $d(p, q)$.

This phenomenon is independent of the nature of the metric space (vectorial or not, in particular) and gives us a way to quantify how hard it is to search on an arbitrary metric space.

Definition 9. The *intrinsic dimensionality* of a metric space is defined as $\rho = \mu^2/2\sigma^2$, where μ and σ^2 are the mean and variance of its histogram of distances.

The technical convenience of the exact definition is made clear shortly. The important part is that the intrinsic dimensionality grows with the mean and decreases with the variance of the histogram.

The particular cases of the L_s distances in vector spaces are useful illustrations. As shown in Yianilos [1999], a uniformly distributed k-dimensional vector space under the L_s distance has mean $\Theta(k^{1/s})$ and standard deviation $\Theta(k^{1/s - 1/2})$. Therefore its intrinsic dimensionality is $\Theta(k)$ (although the constant is not necessarily 1). So the intuitive concept of dimensionality in vector spaces matches our general concept of intrinsic dimensionality.

7.2. A Lower Bound for Pivoting Algorithms

Our main result in this section relates the intrinsic dimensionality with the difficulty of searching with a given search radius r using a pivoting equivalence relation that chooses the pivots at random. As we show next, the difficulty of the problem is related to r and the intrinsic dimensionality ρ.

We are considering independent identically distributed random variables for the distribution of distances between points. Although not accurate, this simplification is optimistic and hence can be used to lower bound the performance of the indexing algorithms. We come back to this shortly.

Let $(q, r)_d$ be a range query over a metric space indexed by means of k random pivots, and let u be an element of $\mathbb{U}$. The probability that u cannot be excluded from direct verification after considering the k pivots is exactly

(a) $Pr(|d(q, p_1) - d(u, p_1)|$
$\leq r, \ldots, |d(q, p_k) - d(u, p_k)| \leq r)$.

Since all the pivots are assumed to be random and their distance distributions independent identically distributed random variables, this expression is the product of probabilities

(b) $Pr(|d(q, p_1) - d(u, p_1)| \leq r) \times \cdots$
$\times Pr(|d(q, p_k) - d(u, p_k)| \leq r)$

which for the same reason can be simplified to

$$Pr(\text{not discarding } u)$$
$$= Pr(|d(q, p) - d(u, p)| \leq r)^k.$$

for a random pivot p.

If X and Y are two independent identically distributed random variables with mean μ and variance σ^2, then the mean of $X - Y$ is 0 and its variance is $2\sigma^2$. Using Chebyschev's inequality[6] we have that $Pr(|X - Y| > \varepsilon) < 2\sigma^2/\varepsilon^2$. Therefore,

$$Pr(|d(q, p) - d(u, p)| \leq r) \geq 1 - \frac{2\sigma^2}{r^2},$$

where σ^2 is precisely the variance of the distance distribution in our metric space. The argument that follows is valid for $2\sigma^2/r^2 < 1$, or $r > \sqrt{2}\sigma$ (large enough radii); otherwise the lower bound is zero. Then, we have

$$Pr(\text{not discarding } u) \geq \left(1 - \frac{2\sigma^2}{r^2}\right)^k.$$

We have now that the total search cost is the number of internal distance

---

[6] For an arbitrary distribution Z with mean μ_z and variance σ_z^2, $Pr(|Z - \mu_z| > \varepsilon) < \sigma_z^2/\varepsilon^2$.

evaluations (k) plus the external evaluations, whose number is on average $n \times Pr$(not discarding u). Therefore

$$\text{Cost} \geq k + n \ \left(1 - \frac{2\sigma^2}{r^2}\right)^k$$

is a *lower bound* to the average search cost by using random pivots. Optimizing we obtain that the best k is

$$k^* = \frac{\ln n + \ln \ln(1/t)}{\ln(1/t)},$$

where $t = 1 - 2\sigma^2/r^2$. Using this optimal k^*, we obtain an absolute (i.e., independent on k) lower bound for the average cost of any random pivot-based algorithm:

$$\begin{aligned}
\text{Cost} \ &\geq \ \frac{\ln n + \ln \ln(1/t) + 1}{\ln(1/t)} \\
&\geq \ \frac{\ln n}{\ln(1/t)} \geq \frac{r^2}{2\sigma^2} \ln n
\end{aligned}$$

which shows that the cost depends strongly on σ/r. As r increases t tends to 1 and the scheme requires more and more pivots and it is anyway more and more costly.

A nice way to represent this result is to assume that we are interested in retrieving a fixed fraction of at most f of the elements, in which case r can be written as $r = \mu - \sigma/\sqrt{f}$ by Chebyschev's inequality. In this case the lower bound becomes

$$\begin{aligned}
\frac{r^2}{2\sigma^2} \ \ln n \ &= \ \frac{(\mu - \sigma/\sqrt{f})^2}{2\sigma^2} \ \ln n \\
&= \ \rho \left(1 - \frac{1}{\sqrt{2\rho f}}\right)^2 \ \ln n \\
&= \ \left(\sqrt{\rho} - \frac{1}{\sqrt{2f}}\right)^2 \ \ln n
\end{aligned}$$

which is valid for $f \ \geq \ 1/(2\rho)$. We have just proved the following.

THEOREM 1. *Any pivot-based algorithm using random pivots has a lower bound $(\sqrt{\rho} - 1/\sqrt{2f})^2 \ln n$ in the average number of distance evaluations performed*

for a random range query retrieving at most a fraction f of the set, where ρ is the intrinsic dimension of the metric space.

This result matches that of Baeza-Yates [1997] and Baeza-Yates and Navorro [1998] on FHQTs, about obtaining $\Theta(\log n)$ search time using $\Theta(\log n)$ pivots, but here we are more interested in the "constant" term, which depends on the dimension.

The theorem shows clearly that the parameters governing the performance of range-searching algorithms are ρ and f. One can expect that f keeps constant as ρ grows, but it is possible that in some applications the query q is known to be a perturbation of some element of $\mathbb{U}$ and therefore we can keep a constant search radius r as the dimension grows. Even in those cases the lower bound may grow if σ shrinks with the dimension, as in the L_s vector spaces with $s > 2$.

We have considered independent identically distributed random variables for each pivot and the query. This is a reasonable approximation, as we do not expect much difference between the "view points" from the general distribution of distances to the individual distributions (see Section 7.4). The expression given in Equation (7.2b) cannot be obtained without this simplification.

A stronger assumption comes from considering all the variables as independent. This is an optimistic consideration equivalent to assuming that in order to discard each element u of the set we take k new pivots at random. The real algorithm fixes k random pivots and uses them to try to discard all the elements u of the set. The latter alternative can suffer from dependencies from a point u to another, which cannot happen in the former case (e.g., if u is close to the third pivot and u' is close to u then the distance from u' to the third pivot carries less information). Since the assumption is optimistic, using it to reduce the joint distribution in Equation (7.2a) to the expression given in Equation (7.2b) keeps the lower bound valid.

Figures 15 and 16 show an experiment on the search cost in $(\mathbb{R}^\ell, L_2)$ using a different number of pivots k and dimensions ℓ.
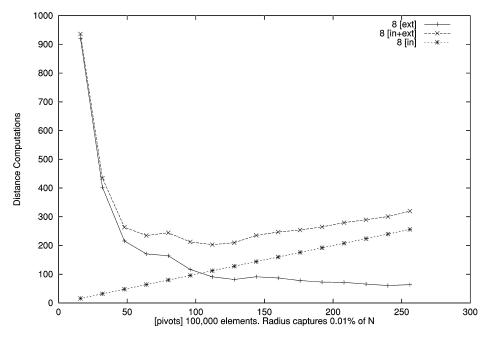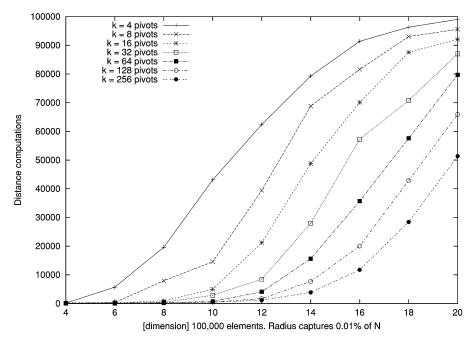
Fig. 15. Internal, external, and overall distance evaluations in eight dimensions, using different numbers of pivots k.

Fig. 16. Overall distance evaluations as the dimension grows for fixed k.

The $n = 100,000$ elements are generated at random and the pivots are randomly chosen from the set. We average over 1,000 random queries whose radius is set to retrieve 10 elements of the set. We count the number of distance evaluations. Figure 15 shows the existence of an optimum $k^* = 110$, and Figure 16 shows the predicted $O(n(1 - 1/\Theta(\ell))^k)$ behavior. We have not enough memory in our machine to show the predicted growth in $k^* \approx \Theta(\ell) \ln(n)$.

Figure 17 shows the effect in a different way. As the dimension grows, the histogram of L_2 moves to the right ($\mu = \Theta(\sqrt{\ell})$). Yet the pivot distance D (in the projected space $(\mathbb{R}^k, L_\infty)$) remains about the same for fixed k. Increasing k from 32 to 512 moves the histogram slightly to the right. This shift is effective in low-dimensional metric spaces, but it is uneffective in high dimensions. The plots of these two histograms can measure how good are the pivoting algorithms. Intuitively, the overlap between the histogram for the pivot distance and the histogram for the original distance is directly proportional to the discriminative power of the pivot mapping. As the overlap increases the algorithms become more effective.

The particular behavior of D in Figure 17 is due to the fact that D is the maximum of k random variables whose *mean* is $\Theta(\sigma)$ (i.e., $|d(p,q) - d(p,x)|$). The fact that D does not grow with ℓ means that, as a lower bound for d, it gets less effective in higher dimensions.

7.3. A Lower Bound for Compact Partitioning Algorithms

We now try to obtain a lower bound to the search cost of algorithms based on the compact partition. The result is surprisingly similar to that of the previous section. Our lower bound considers only the hyperplane criterion, which is the most purely associated with the compact partition. We assume just the same facts about the distribution of distances as in the previous section: all of them are independent identically distributed random variables.

Let $(q, r)_d$ be a range query over a metric space indexed by means of m ran-

dom centers $\{c_1, \ldots, c_m\}$. The m distances $d(q, c_i)$ can be considered as random variables $X_1, \ldots, X_m$ whose distribution is that of the histogram of the metric space. The distribution of the distance from q to its closest center c is that of $Y = \min\{X_1, \ldots, X_m\}$. The hyperplane criterion specifies that a class $[c_i]$ *cannot* be excluded if $d(q, c) + r \geq d(q, c_i) - r$. The probability that this happens is $Pr(Y \geq X_i - 2r)$. But since Y is the minimum over m variables with the same distribution, the probability is $Pr(Z \geq X - 2r)^m$, where X and Z are two random variables distributed according to the histogram. Using Chebyschev's inequality and noticing that if $Z < X - 2r$ then X or Z are at distance at least r from their mean, we can say that
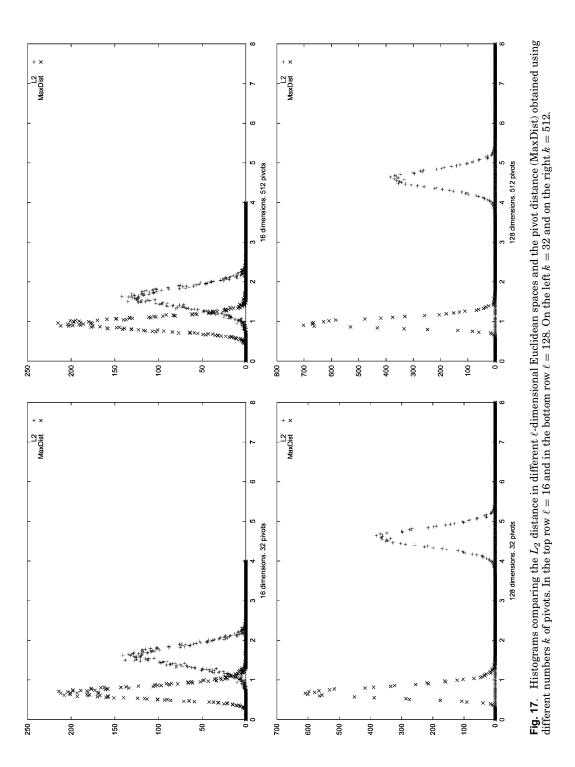
$$Pr(\text{not discarding } [c_i])$$
$$= Pr(Z \geq X - 2r)^m \geq \left(1 - \frac{\sigma^2}{r^2}\right)^m.$$

On average each class has n/m elements, so that the external complexity is $n \times Pr$ (not discarding $[c_i]$). The internal cost to find the intersected classes deserves some discussion. In all the hierarchical schemes that exist, we consider that the real partition is that induced by the leaves of the trees, that is, the most refined ones. We see all the rest of the hierarchy as a mechanism to reduce the internal complexity of finding the small classes (hence the m we use here is not, say, the m of GNATs, but the total number of final classes). It is difficult to determine this internal complexity (an upper bound is m), so we call it $\mathcal{C}_I(m)$, knowing that it is between $\Omega(\log m)$ and $O(m)$. Then a lower bound to the search complexity is

$$\text{Cost} \geq \mathcal{C}_I(m) + n \left(1 - \frac{\sigma^2}{r^2}\right)^m$$

which indeed is very similar to the lower bound on pivot-based algorithms. Optimizing on m yields

$$m^* = \frac{\ln n + \ln \ln(1/t') - \ln \mathcal{C}_I'(m^*)}{\ln(1/t')},$$

Fig. 17. Histograms comparing the L_2 distance in different ℓ-dimensional Euclidean spaces and the pivot distance (MaxDist) obtained using different numbers k of pivots. In the top row $\ell = 16$ and in the bottom row $\ell = 128$. On the left $k = 32$ and on the right $k = 512$.
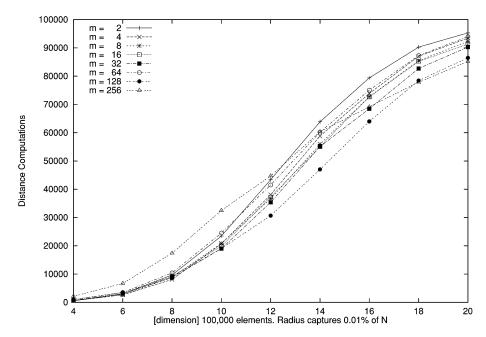
Fig. 18. Overall distance evaluations using a hierarchical compact partitioning with different arities.

where $t' = 1 - \sigma^2/r^2$. Using the optimal m^* the search cost is lower bounded by

$$\text{Cost} = \Omega(\mathcal{C}_I(\log_{1/t'} n)) = \Omega\left(\mathcal{C}_I\left(\frac{r^2}{\sigma^2} \ln n\right)\right)$$

which also shows an increase in the cost as the dimensionality grows. As before we can write $r = \mu - \sigma/\sqrt{f}$. We have just proved the following.

THEOREM 2. *Any compact partitioning algorithm based on random centers has a lower bound $\mathcal{C}_I(2(\sqrt{\rho} - 1/\sqrt{2f})^2)$ in the average number of distance evaluations performed for a random range query retrieving a fraction of at most f of the database, where ρ is the intrinsic dimension of the space and $\mathcal{C}_I()$ is the internal complexity to find the relevant classes, satisfying $\Omega(\log m) = \mathcal{C}_I(m) = O(m)$.*

This result is weaker than Theorem 1 because of our inability to give a good lower bound on $\mathcal{C}_I$, so we cannot ensure more than a logarithmic increase with respect to ρ. However, even assuming $\mathcal{C}_I(m) = \Theta(m)$ (i.e., exhaustive search of the classes), when the theorem becomes very similar to Theorem 1, there is an important reason that explains why the compact partitioning algorithms can in practice be better than pivot-based ones. We can in this case achieve the optimal number of centers m^*, which is impossible in practice for pivot-based algorithms. The reason is that it is much more economical to represent the compact partition using m centers than the pivot partition using k pivots.

Figure 18 shows an experiment on the same dataset, where we have used different m values and a hierarchical compact partitioning based on them. We have used the hyperplane and the covering radius criteria to prune the search. As can be seen, the dependency on the dimension of the space is not so sharp as for pivot-based algorithms, and is closer to a dependency of the form $\Theta(\ell)$.

The general conclusion is that, even if the lower bounds using pivot-based or compact partitioning algorithms look similar, the first ones need much more space to store the classes resulting from k pivots than the last ones using the same

number of partitions. Hence, the latter can realistically use the optimal number of classes, whereas the former cannot. If pivot-based algorithms are given all the necessary memory, then using the optimal number of pivots they can improve over compact partitioning algorithms, because t is better than t', but this is more and more difficult as the dimension grows.

Figure 19 compares both types of partitioning. As can be seen, the pivoting algorithm improves over the compact partitioning if we give it enough pivots. However, "enough" is a number that increases with the dimension and with the fraction retrieved (i.e., ρ and f). For ρ and f large enough, the required number of pivots will be unacceptably high in terms of memory requirements.

7.4. Pivot and Center Selection Techniques

In Faragó et al. [1993], they prove formally that if the dimension is constant, then after properly selecting (not at random!) a constant number k of pivots the exhaustive search costs $O(1)$. This contrasts with our $\Omega(\log n)$ lower bound. The difference is that they do not take the pivots at random but select a set of pivots that is assumed to have certain selectivity properties. This shows that the way in which the pivots are selected can affect the performance. Unfortunately, their pivot selection technique works for vector spaces only.

Little is known about pivot/center (let us call them collectively, "references") selection policies, and in practice most methods choose them at random, with a few exceptions. For instance, in Shapiro [1977] it is recommended to select pivots outside the clusters, and Baeza-Yates et al. [1994] suggest using one pivot from each cluster. All authors agree in that the references should be far apart from each other, which is evident since close references will give almost the same information. On the other hand, references selected at random are already far apart in a high-dimensional space.

The histogram of distances gives a formal characterization of good references. Let us start with a definition.

Definition 10. The local histogram of an element u is the distribution of distances from u to every $x \in \mathbb{X}$.
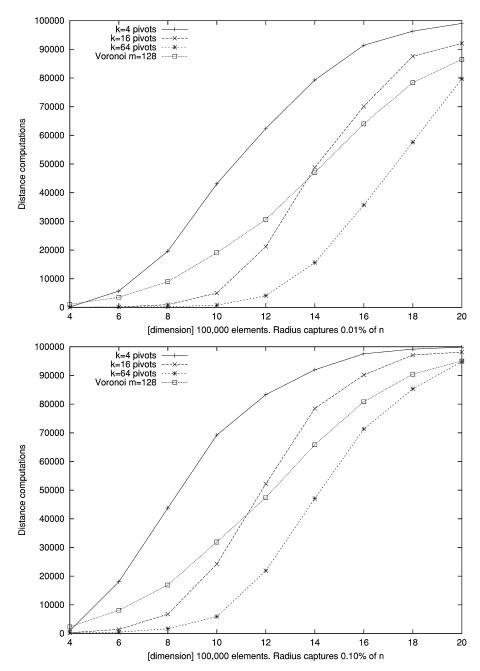
A good reference has a flatter histogram, which means that it will discard more elements at query time. The measure $\rho = \mu^2/(2\sigma^2)$ of intrinsic dimensionality (now defined on the local histogram of u) can be used as a good parameter to evaluate how good a reference is (good references have local histograms with small ρ).

This is related to the difference in viewpoints (histograms) between different references, a subject discussed in depth in Ciaccia et al. [1998a]. The idea is that the local histogram of a reference u may be quite different from the global histogram (especially if u is not selected at random). This is used in Ciaccia et al. [1998a] to show that if the histograms for different references are similar then they can accurately predict the behavior of instances of their data structure (the MT), a completely different issue.

Note also that good reference selection becomes harder as the intrinsic dimension grows. As the global histogram reduces its variance, it becomes more difficult to find references that deviate significantly from it, as already noted in Ciaccia et al. [1998a].

The histogram characterization explains a well-known phenomenon: to discriminate among the elements in the class of a local partition (or in a cluster), it is a good idea to select a pivot from the same class. This makes it more probable to select an element close to them (the ideal would be a centroid). In this case, the distances tend to be smaller and the histogram is not so concentrated in large values. For instance, for LAESA, Micó et al. [1994] do not use the pivots in a fixed order, but the next one is that with minimal L_1 distance to the current candidates. On the other hand, outliers can be good at a first approximation, in order to discriminate among clusters, but later they are unable to separate well the elements of the same cluster.

Selecting good individual references, that is, with a flatter histogram, is not

Fig. 19.   Distance evaluations for increasing dimension. We compare the compact partitioning algorithm of Figure 18 using 128 centers per level against the filtration using k pivots for $k = 4, 16, 64$. On top the search radius captures 0.01% of the set, on the bottom 0.1%.

enough to ensure a good performance. For example, a reference close to a good one is probably good as well, but using both of them gives almost the same information as using only one. It is necessary to include a variety of independent viewpoints in the selection. If we consider the selection of references as an optimization problem, the goal is to obtain a set of references that are collectively good for indexing purposes. The set of references is good if it approximates the original distance well. This is more evident in the case of pivoting algorithms, where a contractive distance is implicitly used. A practical way to decide whether our set of references is effective is to compare the histograms of the original distance and the pivot distance as in Figure 17. A good reference set will give a pivot distance histogram with a large intersection with the histogram of the original distance.

7.5. The Effect on the Discriminative Power

Another reason that explains the curse of dimensionality is related to the decrease in the discriminative power of a partition, because of the odd "shapes" of the classes.

As explained before, a nonlocal partition suffers from the problem of being unable to discriminate between points that are actually far away from each other, which leads to unnecessarily high external complexity. The solution is to select enough pivots or to use a compact partitioning method that yields a local partition. However, even in a local partition, the shape of the cell is fixed at indexing time, while the shape of the space region defined as interesting for a query q is a ball dynamically determined at query time. The discriminative power can be visualized as the volume of the query ball divided by the total volume of the classes intersected by that ball.

A good example is that of a ball inside a box in a k-dimensional L_2 space. The box is the class obtained after using k orthogonal pivots, and the partition obtained is quite local. Yet the volume of the ball is smaller, and the ratio with respect to the volume of the box (i.e., the discriminative power) decreases exponentially with k.

This means that we have to examine a volume (and a number of candidate elements) that, with respect to the size of the final result, grows exponentially with the dimension k. This fact is behind the exponential dependency on the dimension in data structures for vector spaces such as the kd-tree or the R-tree.

The same phenomenon occurs in general metric spaces, where the "shapes" of the cells cannot possibly fit an unknown query ball. We can add more pivots to better bound any possible ball, but this increases the internal complexity and becomes more difficult as the intrinsic dimension grows (the smaller the variance, the more pivots are needed to make a difference).

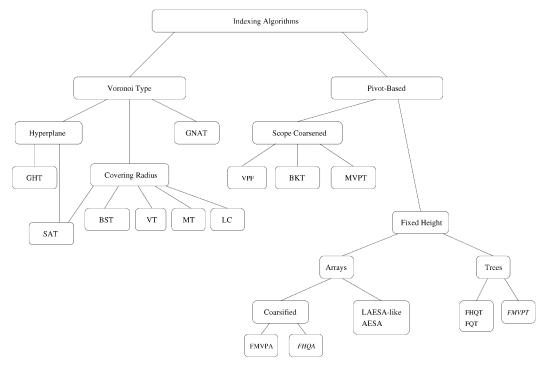## 8. A TAXONOMY OF SEARCH ALGORITHMS

In this section we apply our unifying model to organize all the known approaches in a taxonomy. This helps to identify the essential features of all the existing techniques, to find possible combinations of algorithms not noticed up to now, and to detect which are the most promising areas for optimization.

We first concentrate on pivoting algorithms. They differ in their method to select the pivots, in when the selection is made, and in how much information on the comparisons is used. Later, we consider the compact partitioning algorithms. These differ in their methods to select the centers and to bound the equivalence classes.

Figure 20 summarizes our classification of methods attending to their most important features (explained throughout the section).

8.1. Pivot-Based Algorithms

8.1.1. Search Algorithms. Once we have determined the equivalence relation to use (i.e., the k pivots), we preprocess the dictionary by storing, for each element of $\mathbb{U}$, its k coordinates (i.e., distance to the k pivots). This takes $O(kn)$ preprocessing time and space overhead. The "index" can

Indexing Algorithms

Voronoi Type — Pivot-Based

Hyperplane — GNAT
Scope Coarsened

GHT — Covering Radius
VPF — BKT — MVPT

SAT
BST — VT — MT — LC

Fixed Height

Arrays — Trees

Coarsified — LAESA-like AESA
FHQT FQT — *FMVPT*

FMVPA — *FHQA*

Fig. 20. Taxonomy of the existing algorithms. The methods in italics are combinations that appear naturally as soon as the taxonomy is built.

be seen as a table of n rows and k columns storing $d(u_i, p_j)$.

At query time, we first compare the query q against the k pivots, hence obtaining its k coordinates $[q] = (y_1, \ldots, y_k)$ in the target space, that is, its equivalence class. The cost of this is k evaluations of the distance function d, which corresponds to the internal complexity of the search. We have now to determine, in the target space, which classes may be relevant to the query (i.e., which ones are at distance r or less in the L_∞ metric, which corresponds to the D distance). This does not use further evaluations of d, but it may take extra CPU cost. Finally, the elements belonging to the qualifying classes (i.e., those that cannot be discarded after considering the k pivots) are directly compared against q (external complexity).

The simplest search algorithm proceeds rowwise: consider each element of the set (i.e., each row $(x_1, \ldots, x_k)$ of the table) and see if the triangle inequality allows discarding that row, that is, whether $\max_{i=1..k}\{|x_i - y_i|\} > r$. For each row not discarded using this rule, compare the element directly against q. This is equivalent to traversing the quotient space, using D to discard uninteresting classes.

Although this traversal does not perform more evaluations of d than necessary, it is not the best choice. The reasons are made clear later, as we discover the advantages of alternative approaches. First, notice that the amount of CPU work is $O(kn)$ in the worst case. However, as we abandon a row as soon as we find a difference larger than r along a coordinate, the average case is much closer to $O(n)$ for queries of reasonable selectivity.

The first improvement is to process the set columnwise. That is, we compare the query against the first pivot p_1. Now, we consider the first column of the table and discard all the elements that satisfy $|x_1 - y_1| > r$. Then, we consider the second pivot p_2 and repeat the process only on the elements not discarded up to now. An algorithm implementing this idea is LAESA.

It is not hard to see that the amount of evaluations of d and the total CPU work remains the same as for the rowwise case. However, we can do better now, since each column can be sorted so that the range of qualifying rows can be binary instead of sequentially searched [Nene and Nayar 1997; Chávez et al. 1999]. This is possible because we are interested, at column i, in the values $[y_i - r, y_i + r]$. The extra CPU cost gets closer to $O(k \log n)$ than to $O(n)$ by using this technique.

This is not the only improvement allowed by a columnwise evaluation that cannot be done rowwise. A very important one is that it is not necessary to consider all the k coordinates (recall that we have to perform one evaluation of d to obtain each new query coordinate y_i). As soon as the remaining set of candidates is small enough, we can stop considering the remaining coordinates and directly verify the candidates using the d distance. This point is difficult to estimate beforehand: despite the (few) theoretical results existing [Faragó et al. 1993; Baeza-Yates 1997; Baeza-Yates and Navarro 1998], one cannot normally understand the application well enough to predict the actual optimal number of pivots k^* (i.e., the point where it is better to switch to exhaustive evaluation).

Another improvement that can be done with columnwise evaluation is that the selection of the pivots can be done on the fly instead of beforehand as we have presented it. That is, once we have selected the first pivot p_1 and discarded all the uninteresting elements, the second pivot p_2 may depend on which was the result of p_1. However, for each potential pivot p we have to store the coordinates of all the elements of the set for p (or at least some, as we show later). That is, we select k potential pivots and precompute the table as before, but we can choose in which order the pivots are used (according to the current state of the search) and where we stop using pivots and compare directly.

An extreme case of this idea is AESA, where $k = n$ (i.e., all the elements are potential pivots), and the new pivot at each iteration is selected among the remaining elements. Despite its practical inapplicability because of its $O(n^2)$ preprocessing time and space overhead (i.e., all the distances among the known elements are precomputed), the algorithm performs a surprisingly low number of distance evaluations, much better than when the pivots are fixed. This shows that it is a good idea to select pivots from the current set of candidates (as discussed in the previous sections).

Finally, we notice that instead of a sequential search in the mapped space, we could use an algorithm to search in vector spaces of k dimensions (e.g., kd-trees or R-trees). Depending on their ability to handle larger k values, we could be able to use more pivots without significantly increasing the extra CPU cost. Recall also that, as more pivots are used, the search structures for vector spaces perform worse. This is a very interesting subject which has not been pursued yet, that accounts for balancing between distance evaluations and CPU time.

8.1.2. Coarsening the Equivalence Relation. The alternative of not considering all the k pivots if the remaining set of candidates is small is an example of coarsening an equivalence relation. That is, if we do not consider a given pivot p, we are merging all the classes that differ only in that coordinate. In this case we prefer to coarsen the pivot equivalence relation because computing it with more precision is worse than checking it as is.

There are many other ways to coarsen the equivalence relation, and we cover them here. However, in these cases the coarsening is not done for the sake of reducing the number of distance evaluations, but to improve space usage and precomputation time, as $O(kn)$ can be prohibitively expensive for some applications. Another reason is that, via coarsening, we obtain search algorithms that are sublinear in their extra CPU time. We consider in this section *range coarsening*, *bucket coarsening*, and *scope coarsening*. Their ideas are roughly illustrated in Figure 21.

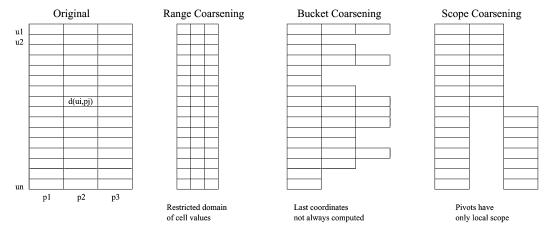It must be clear that all these types of coarsenings *reduce* the discriminative

Original	Range Coarsening	Bucket Coarsening	Scope Coarsening

u1
u2

d(ui,pj)

un

p1   p2   p3

Restricted domain
of cell values

Last coordinates
not always computed

Pivots have
only local scope

Fig. 21.   Different coarsification methods.

power of the resulting equivalence classes, making it necessary to exhaustively consider *more* elements than in the un-coarsened versions of the relations. In the example of the previous section this is amortized by the lower cost to obtain the coarsened equivalence relation. Here we reduce the effectiveness of the algorithms via coarsening, for the sake of reduced preprocessing time and space overhead.

However, space reduction may have a counterpart in time efficiency. If we use less space, then with the same amount of memory we can have more pivots (i.e., larger k). This can result in an overall improvement. The fair way to compare two algorithms is to give them the same amount of memory to use.

8.1.3. Range Coarsening. The auxiliary data structures proposed by most authors for continuous distance functions are aimed at reducing the amount of space needed to store the coordinates of the elements in the mapped space, as well as the time to find the relevant classes. The most popular form is to reduce the precision of d. This is written as

$$x \sim_{p,\{r_i\}} y \Leftrightarrow \exists i, r_i \leq d(x, p) < r_{i+1} \quad \text{and} \\ r_i \leq d(y, p) < r_{i+1}$$

with $\{r_i\}$ a partition of the interval $[0, \infty)$. That is, we assign the same equivalence class to elements falling in the same range of distances with respect to the same pivot

p. This is obviously a coarsening of the previous relation $\sim_p$ and can be naturally extended to more than one pivot.

Figure 2 exemplifies a pivoting equivalence relation where range coarsening is applied, for one pivot. Points in the same ring are in the same equivalence class, despite the fact that their exact distance to the pivot may be different.

A number of actual algorithms use one or another form of this coarsening technique. VPTs and MVPTs divide the distances in slices so that the same number of elements lie in each slice (note that the slices are different for each pivot). VPTs use two slices and MVPTs use many. Their goal is to obtain balanced trees. BKTs, FQTs, and FHQTs, on the other hand, propose range coarsening for continuous distance functions but do not specify how to coarsen.

In this work we consider that the "natural" extension of BKTs, FQTs, and FHQTs assigns slices of the same width to each branch, and that the tree has the same arity across all its nodes. At each node, the slice width is recomputed so that using slices of that fixed width the node has the desired arity.

Therefore, we can have slices of fixed width (BKT, FQT, FHQT) or determined by percentiles (VPT, MVPT, FQA). We may have a different pivot per node (BKT, VPT, MVPT) or per level (FQT, FHQT, FQA). Among the last, we can define the

Table II. Different Options for Range Coarsening

		Fixed Percentiles	Fixed Width
Different pivot per node (scope coarsening)		VPT, MVPT	BKT
Different pivot per level	Different slice per node	*FMVPT*	FQT, FHQT
	Different slice per level	FMVPA (FQA)	*FHQA*

[a]We put in italics the new structures created to fill the empty holes.

slices at each node (FQT, FHQT) or for the whole level (FQA). All these choices are drawn in Table II. The empty slots have been filled with new structures that are now defined.

8.1.3.1. FHQA. This is similar to an FQA except that the slices are of fixed width. At each *level* the slice width is recomputed so that a maximum arity is guaranteed. In the FHQT, instead, each *node* has a different slice width.

8.1.3.2. FMVPT. This is a cross between an MVPT and a FHQT. The range of values is divided using the $m-1$ uniform percentiles to balance the tree, as in MVPTs. The tree has a fixed height h, as FHQTs. At each node the ranges are recomputed according to the elements lying in that subtree. The particular case where $m = 2$ is called FHVPT.

8.1.3.3. FMVPA. This is just a new name for the FQA, more appropriate for our discussion since it is to MVPTs as FHQAs are to FHQTs: the FMVPA uses variable width slices to ensure that the subtrees are balanced in size, but the same slices are used for all the nodes of a single level, so that the balance is only levelwise.

The combinations we have just created allow us to explain some important concepts.
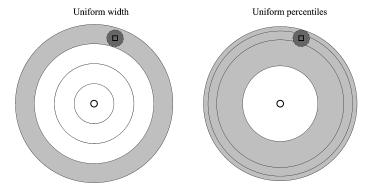
8.1.3.4. Amount of Range Coarsening. Let us consider FHQAs and FMVPAs. They are no more than LAESA with different forms of range coarsening. They use k fixed pivots and use b bits to represent the coordinates (i.e., the distances from each point to each of the h pivots). So only 2^b different values can be expressed. The two structures differ only in how they

coarsen the distances to put them into 2^b ranges. Their total space requirement is then reduced to kbn bits.

However, range coarsening is not just a technique to reduce space, but the same space can be used to accommodate more pivots. It is not immediate how convenient it is to coarsen in order to use more pivots, but it is clear that this technique can improve the overall effectiveness of the algorithm.

8.1.3.5. Percentiles Versus Fixed Width. Another unclear issue is whether fixed slices are better or worse than percentile splitting. A balanced data structure has obvious advantages because the internal complexity may be reduced. Fixed slices produce unbalanced structures since the outer rings have many more elements (especially on high dimensions). On the other hand, in high dimensions the outer rings tend to be too narrow if percentile splitting is used (because a small increment in radius gets many new elements inside the ring). If the rings are too narrow, many rings will be frequently included in the radius of interest of the queries (see Figure 22). An alternative idea is shown in Chávez [1999], where the slices are optimized to minimize the number of branches that must be considered. In this case, each class can be an arbitrary set of slices.

8.1.3.6. Trees Versus Arrays. FHQTs and FMVPTs are almost tree versions of FHQAs and FMVPAs, respectively. They are m-ary trees where all the elements belonging to the same coarsened class are stored in the same subtree. Instead of explicitly storing the m coordinates, the trees store them *implicitly*: the elements are at the leaves, and their paths from the

Uniform width   Uniform percentiles

Fig. 22. The same query intersects more rings when using uniform percentiles.

root spell out the coarsened coordinate values. This makes the space requirements closer to $O(n)$ in practice, instead of $O(bkn)$ (although the constant is very low for the array versions, which may actually take less space). Moreover, the search for the relevant elements can be organized using the tree: if all the interesting elements have their first coordinate in the ith ring, then we just traverse the ith subtree. This reduces the extra CPU time. If the distance is too fine-grained, however, the root will have nearly n children and the subtrees will have just one child. The structure will be very similar to a table of k coordinates per element and the search will degenerate into a linear row-wise traversal. Hence, range coarsening is also a tool to reduce the extra CPU time.

We have given the trees the ability to define the slices at each node instead of at each level as do the array versions. This allows them to adapt better to the data, but the values of the slices used need more space. It is not clear whether it pays to store all these slice values.

Summarizing, range coarsening can be applied using fixed-width or fixed-percentile slices. They can reduce the space necessary to store the coordinates, which can allow the use of more pivots with the same amount of memory. Therefore, it is not just a technique to reduce space but it can improve the search complexity. Range coarsening can also be used to organize treelike search schemes that are sublinear in extra CPU time.

8.1.4. Bucket Coarsening. To reduce space requirements in the above trees, we can avoid building subtrees that have few elements. Instead, all their elements are stored in a *bucket*. When the search arrives at a bucket, it has to exhaustively consider all the elements.

This is a form of coarsening, since for the elements in the bucket we do not consider the last pivots, and it resembles the previous idea (Section 8.1.1) of not computing the k pivots. However, in this case the decision is taken offline, at index construction time, and this allows reducing space by not storing those coordinates. In the previous case the decision was taken at search time. The crucial difference is that if the decision is taken at search time, we can know exactly the total amount of exhaustive work to do by not taking further coordinates. However, in an offline decision we can only consider the search along this branch of the tree, and we cannot predict how many branches will be considered at search time.

This idea is used for discrete distance functions in FQTs, which are similar to FHQTs except for the use of buckets. It has been also applied to continuous setups to reduce space requirements further.

8.1.5. Scope Coarsening. The last and least obvious form of coarsening is the one we call "scope coarsening." In fact, the use of this form of coarsening makes it difficult to notice that many algorithms based on trees are in fact pivoting algorithms.

This coarsening is based on, instead of storing all the coordinates of all the elements, just storing some of them. Hence, comparing the query against some pivots helps to discriminate on some subset of the database only. To use this fact to reduce space, we must determine offline which elements will store their distance to which pivots. There are many ways to use this idea, but it has been used only in the following way.

In FHVPTs there is a single pivot per level of the tree, as in FHQTs. The left and right subtrees of VPTs, on the other hand, use *different* pivots. That is, if we have to consider both the left and right subtrees (because the radius r does not allow us to completely discard one), then comparing the query q against the left pivot will be useful for the left subtree only. There is no information stored about the distances from the left pivot to the elements of the right subtree, and vice versa. Hence, we have to compare q against both pivots. This continues recursively. The same idea is used for BKTs and MVPTs.

Although at first sight it is clear that we reduce space, this is not the direct way in which the idea is used in those schemes. Instead, they combine it with a huge increase in the number of potential pivots. For each subtree, an element belonging to the subtree is selected as the pivot and deleted from the set. If no bucketing is used, the result is a tree where each element is a node somewhere in the tree and hence a potential pivot. The tree takes $O(n)$ space, which shows that we can successfully combine a large number of pivots with scope coarsening to have low space requirements (n instead of n^2 as in AESA).

The possible advantage (apart from guaranteed linear space and slightly reduced space in practice) of these structures over those that store all the coordinates (as FQTs and FHQTs) is that the pivots are better suited to the searched elements in each subtree, since they are selected from the same subset. This same property makes AESA a good (although impractical) algorithm.

Shapiro [1977] and Bozkaya and Ozsoyoglu [1997] propose hybrids (for BKT and VPT, resp.) where a number of fixed pivots are used at each node, and for each resulting class a new set of pivots is selected. Note that, historically, FQTs and FHQTs are an evolution over BKTs.

8.2. Compact Partitioning Algorithms

All the remaining algorithms (GHTs, BSTs, GNATs, VTs, MTs, SATs, LCs) rely on a hierarchical compact partition of the metric space. A first source of differences is in how the centers are selected at each node. GHTs and BSTs take two elements per level. VTs repeat previous centers when creating new nodes. GNATs select m centers far apart. MTs try to minimize covering radii. SATs select a variable number of close neighbors of the parent node. LC chooses at random one center per cluster.

The main difference, however, lies in the search algorithm. While GHTs use purely the hyperplane criterion, BSTs, VTs, MTs, and LCs use only the covering radius criterion. SATs use both criteria to increase pruning. In all these algorithms the query q is compared against all the centers of the current node and the criteria are used to discard subtrees.

GNATs are a little different, as they use none of the above criteria. Instead, they apply an AESA-like search over the m centers considering their "range" values. That is, a center c_i is selected, and if the query ball does not intersect a ring around c_i that contains all the elements of c_j, then c_j and all its class (subtree) can be safely discarded. In other words, GNATs limit the class of each center by intersecting rings around the other centers. This way of limiting the extent of a class is different from both the hyperplane and the covering radius criteria. It is probably more efficient, but it requires storing $O(m^2)$ distances at each level.

Table III summarizes the differences. It is clear that there are many possible combinations that have not been tried, but we do not attempt to enumerate all of them.

The compact partition is an attempt to obtain local classes, more local than those based on pivots. A general technique to do this is to identify clusters of close objects

Table III.   Different Options for Limiting Classes

With hyperplanes	GHT, SAT
With balls	BST, VT, MT, SAT, LC
With rings	GNAT

in the set. There exist many clustering algorithms to build equivalence relations. However, most are defined on vector spaces instead of general metric spaces. An exception is Brito et al. [1996], who report very good results. However, it is not clear that good clustering algorithms directly translate into good algorithms for proximity searching. Another clustering algorithm, based on cliques, is presented in Burkhard and Keller [1973], but the results are similar to the simpler BKT. This area is largely unexplored, and the developments here could be converted into improved search algorithms.

9. CONCLUSIONS

Metric spaces are becoming a popular model for similarity retrieval in many unrelated areas. We have surveyed the algorithms that index metric spaces to answer proximity queries. We have not just enumerated the existing approaches to discuss their good and bad points. We have, in addition, presented a unified framework that allows understanding the existing approaches under a common view. It turns out that most of the existing algorithms are indeed variations on a few common ideas, and by identifying them, previously unnoticed combinations have naturally appeared. We have also analyzed the main factors that affect the efficiency when searching metric spaces. The main conclusions of our work are summarized as follows.

1. The concept of intrinsic dimensionality can be defined on general metric spaces as an abstract and quantifiable measure that affects the search performance.

2. The main factors that affect the efficiency of the search algorithms are the intrinsic dimensionality of the space and the search radius.

3. Equivalence relations are a common ground underlying all the index-

ing algorithms, and they divide the search cost in terms of internal and external complexity.

4. A large class of search algorithms relies on taking k pivots and mapping the metric space onto $\mathbb{R}^k$ using the L_∞ distance. Another important class uses compact partitioning.

5. The equivalence relations can be coarsened to save space or to improve the overall efficiency by making better use of the pivots. A hierarchical refinement of classes can improve performance.

6. Although there is an optimal number of pivots to use, this number is too high in terms of space requirements. In practical terms, a pivot-based index can outperform a compact partitioning index if it has enough memory.

7. As this amount of memory becomes infeasible as the dimension grows, compact partitioning algorithms normally outperform pivot-based ones in high dimensions.

8. In high dimensions the search radius needed to retrieve a fixed percentage of the database is very large. This is the reason for the failure to overcome the brute force search with an exact indexing algorithm.

A number of open issues require further attention. The main ones follow.

—For pivot-based algorithms, understand better the effect of pivot selection, devising methods to choose effective pivots. The subject of the appropriate number of pivots and its relation to the intrinsic dimensionality of the space plays a role here. The histogram of distances may be a good tool for pivot selection.

—For compact partitioning algorithms, work more on clustering schemes in order to select good centers. Find ways to reduce construction times (which are in many cases too high).

—Search for good hybrids between compact partitioning and pivoting algorithms. The first ones cope better with high dimensions and the second ones improve as more memory is given

to them. After the space is clustered the intrinsic dimension of the clusters is smaller, so a top-level clustering structure joined with a pivoting scheme for the clusters is an interesting alternative. Those pivots should be selected from the cluster because the clusters have high locality.

—Take extra CPU complexity into account, which we have barely considered in this work. In some applications the distance is not so expensive that one can disregard any other type of CPU cost. The use of specialized search structures in the mapped space (especially $\mathbb{R}^k$) and the resulting complexity tradeoff deserves more attention.

—Take I/O costs into account, which may very well dominate the search time in some applications. The only existing work on this is the M-tree [Ciaccia et al. 1997].

—Focus on nearest neighbor search. Most current algorithms for this problem are based on range searching, and despite that the existing heuristics seem difficult to improve, truly independent ways to address the problem could exist.

—Consider approximate and probabilistic algorithms, which may give much better results at a cost that, especially for this problem, seems acceptable. A first promising step is Chávez and Navarro [2001b].

REFERENCES

APERS, P., BLANKEN, H., AND HOUTSMA, M. 1997. *Multimedia Databases in Perspective*. Springer-Verlag, New York.

ARYA, S., MOUNT, D., NETANYAHU, N., SILVERMAN, R., AND WU, A. 1994. An optimal algorithm for approximate nearest neighbor searching in fixed dimension. In *Proceedings of the Fifth ACM–SIAM Symposium on Discrete Algorithms* (*SODA'94*), 573–583.

AURENHAMMER, F. 1991. Voronoi diagrams—A survey of a fundamental geometric data structure. *ACM Comput. Surv. 23*, 3.

BAEZA-YATES, R. 1997. Searching: An algorithmic tour. In *Encyclopedia of Computer Science and Technology*, vol. 37, A. Kent and J. Williams, Eds., Marcel Dekker, New York, 331–359.

BAEZA-YATES, R. AND NAVARRO, G. 1998. Fast approximate string matching in a dictionary. In *Proceedings of the Fifth South American Symposium on String Processing and Information Retrieval* (*SPIRE'98*), IEEE Computer Science Press, Los Alamitos, Calif., 14–22.

BAEZA-YATES, R. AND RIBEIRO-NETO, B. 1999. *Modern Information Retrieval*. Addison-Wesley, Reading, Mass.

BAEZA-YATES, R., CUNTO, W., MANBER, U., AND WU, S. 1994. Proximity matching using fixed-queries trees. In *Proceedings of the Fifth Combinatorial Pattern Matching* (*CPM'94*), Lecture Notes in Computer Science, vol. 807, 198–212.

BENTLEY, J. 1975. Multidimensional binary search trees used for associative searching. *Commun. ACM 18*, 9, 509–517.

BENTLEY, J. 1979. Multidimensional binary search trees in database applications. *IEEE Trans. Softw. Eng. 5*, 4, 333–340.

BENTLEY, J., WEIDE, B., AND YAO, A. 1980. Optimal expected-time algorithms for closest point problems. *ACM Trans. Math. Softw. 6*, 4, 563–580.

BERCHTOLD, S., KEIM, D., AND KRIEGEL, H. 1996. The X-tree: An index structure for high-dimensional data. In *Proceedings of the 22nd Conference on Very Large Databases* (*VLDB'96*), 28–39.

BHANU, B., PENG, J., AND QING, S. 1998. Learning feature relevance and similarity metrics in image databases. In *Proceedings of the IEEE Workshop on Content-Based Access of Image and Video Libraries* (Santa Barbara, Calif.), IEEE Computer Society, Los Alamitos, Calif., 14–18.

BIMBO, A. D. AND VICARIO, E. 1998. Using weighted spatial relationships in retrieval by visual contents. In *Proceedings of the IEEE Workshop on Content-Based Access of Image and Video Libraries* (Santa Barbara, Calif.), IEEE Computer Society, Los Alamitos, Calif., 35–39.

BÖHM, C., BERCHTOLD, S., AND KEIM, A. 2002. Searching in high-dimensional spaces—index structures for improving the performance of multimedia databases. *ACM Comput. Surv.* To appear.

BOZKAYA, T. AND OZSOYOGLU, M. 1997. Distance-based indexing for high-dimensional metric spaces. In *Proceedings of ACM SIGMOD International Conference on Management of Data*, SIGMOD *Rec. 26*, 2, 357–368.

BRIN, S. 1995. Near neighbor search in large metric spaces. In *Proceedings of the 21st Conference on Very Large Databases* (*VLDB'95*), 574–584.

BRITO, M., CHÁVEZ, E., QUIROZ, A., AND YUKICH, J. 1996. Connectivity of the mutual k-nearest neighbor graph in clustering and outlier detection. *Stat. Probab. Lett. 35*, 33–42.

BUGNION, E., FHEI, S., ROOS, T., WIDMAYER, P., AND WIDMER, F. 1993. A spatial index for approximate multiple string matching. In *Proceedings of the First South American Workshop on String Processing* (*WSP'93*), R. Baeza-Yates and N. Ziviani, Eds., 43–53.

BURKHARD, W. AND KELLER, R. 1973. Some approaches to best-match file searching. *Commun. ACM 16*, 4, 230–236.

CASCIA, M. L., SETHI, S., AND SCLAROFF, S. 1998. Combining textual and visual cues for content-based image retrieval on the world wide web. In *Proceedings of the IEEE Workshop on Content-Based Access of Image and Video Libraries* (Santa Barbara, Calif.), IEEE Computer Society, Los Alamitos, Calif., 24–28.

CHÁVEZ, E. 1999. Optimal discretization for pivot based algorithms. Manuscript. `ftp:// garota.fismat.umich.mx/pub/users/elchavez/ minimax.ps.gz`.

CHÁVEZ, E. AND MARROQUÍN, J. 1997. Proximity queries in metric spaces. In *Proceedings of the Fourth South American Workshop on String Processing* (*WSP'97*), R. Baeza-Yates, Ed. Carleton University Press, Ottawa, 21–36.

CHÁVEZ, E. AND NAVARRO, G. 2000. An effective clustering algorithm to index high dimensional metric spaces. In *Proceedings of the Seventh String Processing and Informational Retrieval* (*SPIRE'00*), IEEE CS Press, 75–86.

CHÁVEZ, E. AND NAVARRO, G. 2001a. Towards measuring the searching complexity of metric spaces. In *Proceedings of the Mexican Computing Meeting*, vol. II, 969–972.

CHÁVEZ, E. AND NAVARRO, G. 2001b. A probabilistic spell for the curse of dimensionality. In *Proceedings of the Third Workshop on Algorithm Engineering and Experimentation* (*ALENEX'01*), 147–160. Lecture Notes in Computer Science v. 2153.

CHÁVEZ, E. AND NAVARRO, G. 2002. A metric index for approximate string matching. In *Proceedings of the Fifth Latin American Symposium on Theoretical Informatics* (*LATIN'02*), To appear, Lecture Notes in Computer Science.

CHÁVEZ, E., MARROQUÍN, J., AND BAEZA-YATES, R. 1999. Spaghettis: An array based algorithm for similarity queries in metric spaces. In *Proceedings of String Processing and Information Retrieval* (*SPIRE'99*), IEEE Computer Science Press, Los Alamitos, Calif., 38–46.

CHÁVEZ, E., MARROQUIN, J. L., AND NAVARRO, G. 2001. Fixed queries array: a fast and economical data structure for proximity searching. *Multimedia Tools and Applications 14* (*2*), 113–135. Kluwer.

CHAZELLE, B. 1994. Computational geometry: A retrospective. In *Proceedings of the 26th ACM Symposium on the Theory of Computing* (*STOC'94*), 75–94.

CHIUEH, T. 1994. Content-based image indexing. In *Proceedings of the Twentieth Conference on Very Large Databases* (*VLDB'94*), 582–593.

CIACCIA, P., PATELLA, M., AND ZEZULA, P. 1997. M-tree: An efficient access method for similarity search in metric spaces. In *Proceedings of the 23rd Conference on Very Large Databases* (*VLDB'97*), 426–435.

CIACCIA, P., PATELLA, M., AND ZEZULA, P. 1998a. A cost model for similarity queries in metric spaces. In *Proceedings of the Seventeenth ACM SIGACT–SIGMOD-SIGART Symposium on Principles of Database Systems* (*PODS'98*).

CIACCIA, P., PATELLA, M., AND ZEZULA, P. 1998b. Processing complex similarity queries with distance-based access methods. In *Proceedings of the Sixth International Conference on Extending Database Technology* (*EDBT'98*).

CLARKSON, K. 1999. Nearest neighbor queries in metric spaces. *Discrete Comput. Geom. 22*, 1, 63–93.

COX, T. AND COX, M. 1994. *Multidimensional Scaling*. Chapman and Hall, London.

DEHNE, F. AND NOLTEMEIER, H. 1987. Voronoi trees and clustering problems. *Inf. Syst. 12*, 2, 171–175.

DEVROYE, L. 1987. *A Course in Density Estimation*. Birkhauser, Boston.

DUDA, R. AND HART, P. 1973. *Pattern Classification and Scene Analysis*. Wiley, New York.

FALOUTSOS, C. AND KAMEL, I. 1994. Beyond uniformity and independence: Analysis of R-trees using the concept of fractal dimension. In *Proceedings of the Thirteenth ACM Symposium on Principles of Database Principles* (*PODS'94*), 4–13.

FALOUTSOS, C. AND LIN, K. 1995. Fastmap: A fast algorithm for indexing, data mining and visualization of traditional and multimedia datasets. *ACM SIGMOD Rec. 24*, 2, 163–174.

FALOUTSOS, C., EQUITZ, W., FLICKNER, M., NIBLACK, W., PETKOVIC, D., AND BARBER, R. 1994. Efficient and effective querying by image content. *J. Intell. Inf. Syst. 3*, 3/4, 231–262.

FARAGÓ, A., LINDER, T., AND LUGOSI, G. 1993. Fast nearest-neighbor search in dissimilarity spaces. *IEEE Trans. Patt. Anal. Mach. Intell. 15*, 9, 957–962.

FRAKES, W. AND BAEZA-YATES, R., EDS. 1992. *Information Retrieval: Data Structures and Algorithms*. Prentice-Hall, Englewood Cliffs, N.J.

GAEDE, V. AND GÜNTHER, O. 1998. Multidimensional access methods. *ACM Comput. Surv. 30*, 2, 170–231.

GUTTMAN, A. 1984. R-trees: A dynamic index structure for spatial searching. In *Proceedings of the ACM SIGMOD International Conference on Management of Data*, 47–57.

HAIR, J., ANDERSON, R., TATHAM, R., AND BLACK, W. 1995. *Multivariate Data Analysis with Readings*, 4th ed. Prentice-Hall, Englewood Cliffs, N.J.

HJALTASON, G. AND SAMET, H. 1995. Ranking in spatial databases. In *Proceedings of Fourth International Symposium on Large Spatial Databases*, 83–95.

JAIN, A. AND DUBES, R. 1988. *Algorithms for Clustering Data*. Prentice-Hall, Englewood Cliffs, N.J.

KALANTARI, I. AND MCDONALD, G. 1983. A data structure and an algorithm for the nearest point problem. *IEEE Trans. Softw. Eng. 9*, 5.

MEHLHORN, K. 1984. *Data Structures and Algorithms*, Volume III—Multidimensional Searching and Computational Geometry. Springer-Verlag, New York.

MICÓ, L., ONCINA, J., AND CARRASCO, R. 1996. A fast branch and bound nearest neighbour classifier in metric spaces. *Patt. Recogn. Lett. 17*, 731–739.

MICÓ, L., ONCINA, J., AND VIDAL, E. 1994. A new version of the nearest-neighbor approximating and eliminating search (AESA) with linear preprocessing-time and memory requirements. *Patt. Recog. Lett. 15*, 9–17.

NAVARRO, G. 1999. Searching in metric spaces by spatial approximation. In *Proceedings of String Processing and Information Retrieval (SPIRE'99)*, IEEE Computer Science Press, Los Alamitos, Calif., 141–148.

NAVARRO, G. AND REYES, N. 2001. Dynamic spatial approximation trees. In *Proceedings of the Eleventh Conference of the Chilean Computer Science Society (SCCC'01)*, IEEE CS Press, 213–222.

NENE, S. AND NAYAR, S. 1997. A simple algorithm for nearest neighbor search in high dimensions. *IEEE Trans. Patt. Anal. Mach. Intell. 19*, 9, 989–1003.

NIEVERGELT, J. AND HINTERBERGER, H. 1984. The grid file: An adaptable, symmetric multikey file structure. *ACM Trans. Database Syst. 9*, 1, 38–71.

NOLTEMEIER, H. 1989. Voronoi trees and applications. In *Proceedings of the International Workshop on Discrete Algorithms and Complexity* (Fukuoka, Japan), 69–74.

NOLTEMEIER, H., VERBARG, K., AND ZIRKELBACH, C. 1992. Monotonous bisector* trees—A tool for efficient partitioning of complex schenes of geometric objects. In *Data Structures and Efficient Algorithms*, Lecture Notes in Computer Science, vol. 594, Springer-Verlag, New York, 186–203.

PRABHAKAR, S., AGRAWAL, D., AND ABBADI, A. E. 1998. Efficient disk allocation for fast similarity searching. In *Proceedings of ACM SPAA'98* (Puerto Vallarta, Mexico).

ROUSSOPOULOS, N., KELLEY, S., AND VINCENT, F. 1995. Nearest neighbor queries. In *Proceedings of the ACM SIGMOD'95*, 71–79.

SALTON, G. AND MCGILL, M. 1983. *Introduction to Modern Information Retrieval*. McGraw-Hill, New York.

SAMET, H. 1984. The quadtree and related hierarchical data structures. *ACM Comput. Surv. 16*, 2, 187–260.

SANKOFF, D. AND KRUSKAL, J., EDS. 1983. *Time Warps, String Edits, and Macromolecules: The Theory and Practice of Sequence Comparison*. Addison-Wesley, Reading, Mass.

SASHA, D. AND WANG, T. 1990. New techniques for best-match retrieval. *ACM Trans. Inf. Syst. 8*, 2, 140–158.

SHAPIRO, M. 1977. The choice of reference points in best-match file searching. *Commun. ACM 20*, 5, 339–343.

SUTTON, R. AND BARTO, A. 1998. *Reinforcement Learning : An Introduction*. MIT Press, Cambridge, Mass.

UHLMANN, J. 1991a. Implementing metric trees to satisfy general proximity/similarity queries. Manuscript.

UHLMANN, J. 1991b. Satisfying general proximity/ similarity queries with metric trees. *Inf. Proc. Lett. 40*, 175–179.

VERBARG, K. 1995. The C-Ttree: A dynamically balanced spatial index. *Comput. Suppl. 10*, 323–340.

VIDAL, E. 1986. An algorithm for finding nearest neighbors in (approximately) constant average time. *Patt. Recogn. Lett. 4*, 145–157.

WATERMAN, M. 1995. *Introduction to Computational Biology*. Chapman and Hall, London.

WEBER, R. SCHEK, H.-J., AND BLOTT, S. 1998. A quantitative analysis and performance study for similarity-search methods in high-dimensional spaces. In *Proceedings of the International Conference on Very Large Databases (VLDB'98)*.

WHITE, D. AND JAIN, R. 1996. Algorithms and strategies for similarity retrieval. Tech. Rep. VCL-96-101 (July), Visual Computing Laboratory, University of California, La Jolla.

YAO, A. 1990. *Computational Geometry*, J. Van Leeuwen, Ed. Elsevier Science, New York, 345–380.

YIANILOS, P. 1993. Data structures and algorithms for nearest neighbor search in general metric spaces. In *Proceedings of the Fourth ACM–SIAM Symposium on Discrete Algorithms (SODA'93)*, 311–321.

YIANILOS, P. 1999. Excluded middle vantage point forests for nearest neighbor search. In *DIMACS Implementation Challenge, ALENEX'99* (Baltimore, Md).

YIANILOS, P. 2000. Locally lifting the curse of dimensionality for nearest neighbor search. In *Proceedings of the Eleventh ACM–SIAM Symposium on Discrete Algorithms (SODA'00)*, 361–370.

YOSHITAKA, A. AND ICHIKAWA, T. 1999. A survey on content-based retrieval for multimedia databases. *IEEE Trans. Knowl. Data Eng. 11*, 1, 81–93.